
Point Operations

Prof. George Wolberg
Dept. of Computer Science
City College of New York

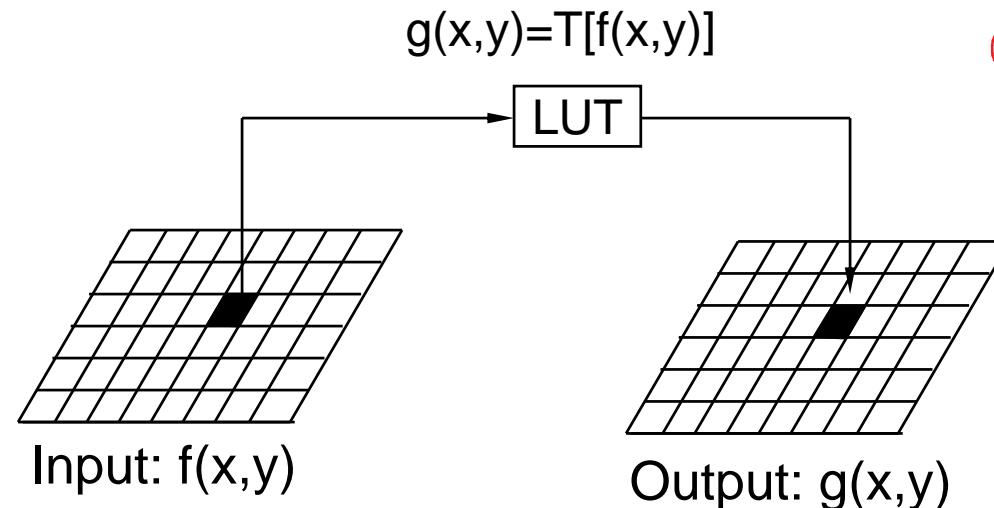
Objectives

- In this lecture we describe point operations commonly used in image processing:
 - Thresholding
 - Quantization (aka posterization)
 - Gamma correction
 - Contrast/brightness manipulation
 - Histogram equalization/matching

Point Operations

lookup
table

- Output pixels are a function of only one input point:
$$g(x,y) = T[f(x,y)]$$
- Transformation T is implemented with a lookup table:
 - An input value indexes into a table and the data stored there is copied to the corresponding output position.
 - The LUT for an 8-bit image has 256 entries.



Point operation
output pixel is a function
output of one and only one
input pixel

Mapping one gray level to another basically.

Graylevel Transformations

- **Point transformation:** changes a pixel's value without changing its location.

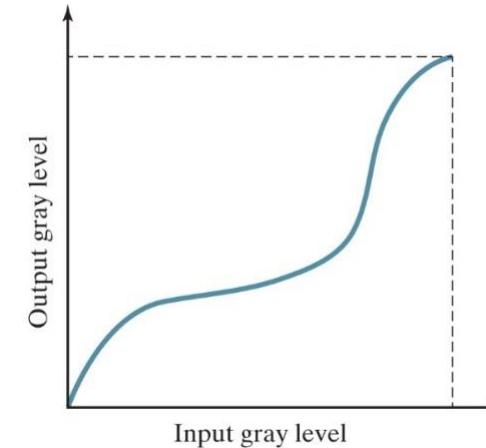
$$I'(x, y) = f(I(x, y))$$

or T

Output image Input image

Figure 3.7 A graylevel transformation maps input gray levels to output gray levels. Based on http://www.unit.eu/cours/videocommunication/Point_Transformation_histogram.pdf

F:



ALGORITHM 3.5 Transform gray levels of an image

```
TRANSFORMGRAYLEVELS( $I, f$ )
```

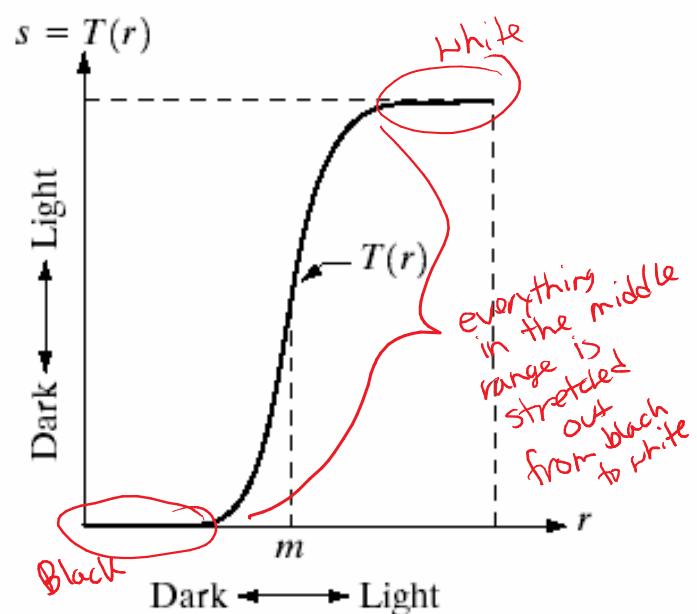
Input: grayscale image I , graylevel mapping f

Output: transformed image I'

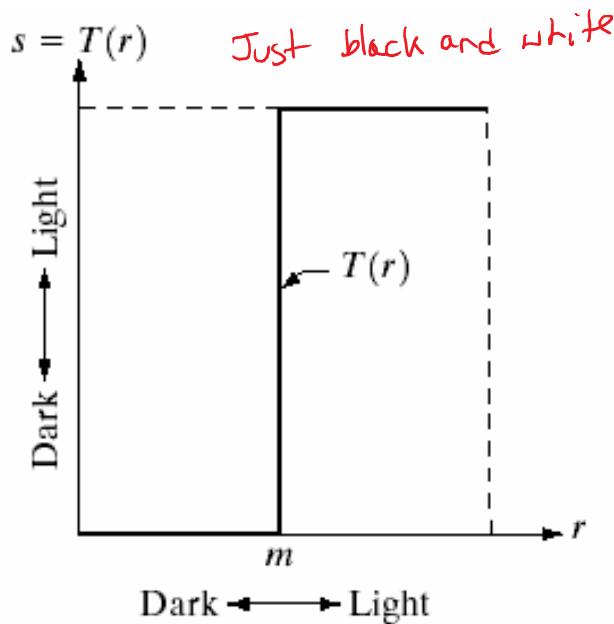
```
1 for  $(x, y) \in I$  do
2      $I'(x, y) \leftarrow f(I(x, y))$ 
3 return  $I'$ 
```

Graylevel Transformation T

Contrast enhancement:
Darkens levels below m
Brightens levels above m



Thresholding:
Replace values below m to black (0)
Replace values above m to white (255)



a | b

FIGURE 3.2 Gray-level transformation functions for contrast enhancement.

We can sample all 256 input values
and put them into a table.

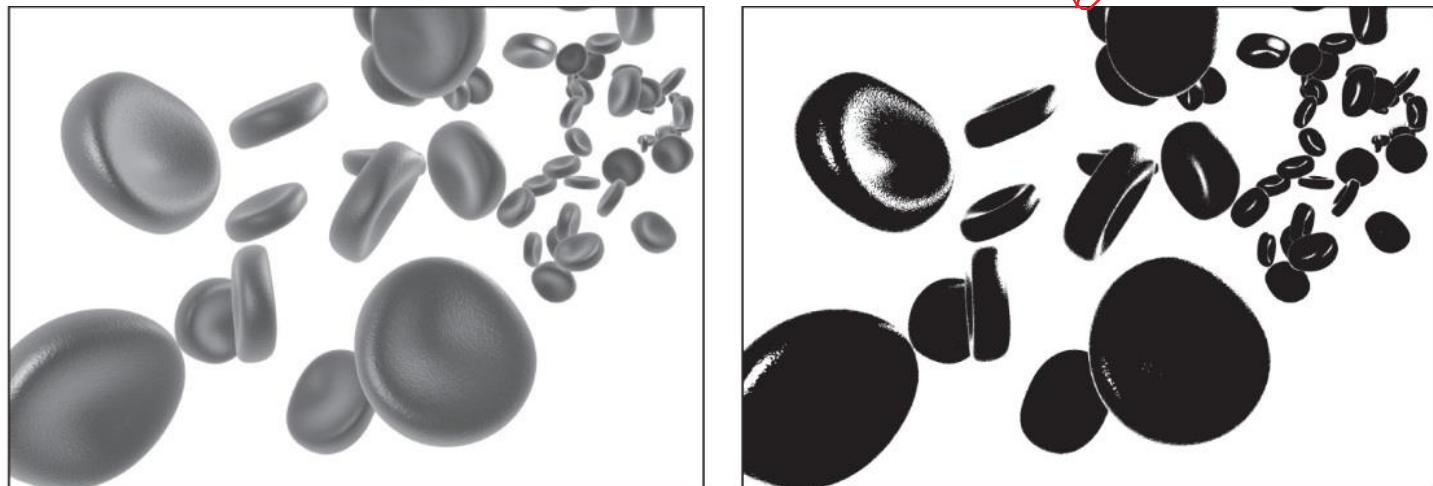
Thresholding

- **Thresholding:** takes a grayscale image and sets every output pixel to 1 if its input gray level is above a certain threshold, or to 0 otherwise:

$$I'(x, y) = \begin{cases} 1 & \text{if } I(x, y) > \tau \\ 0 & \text{otherwise} \end{cases}$$

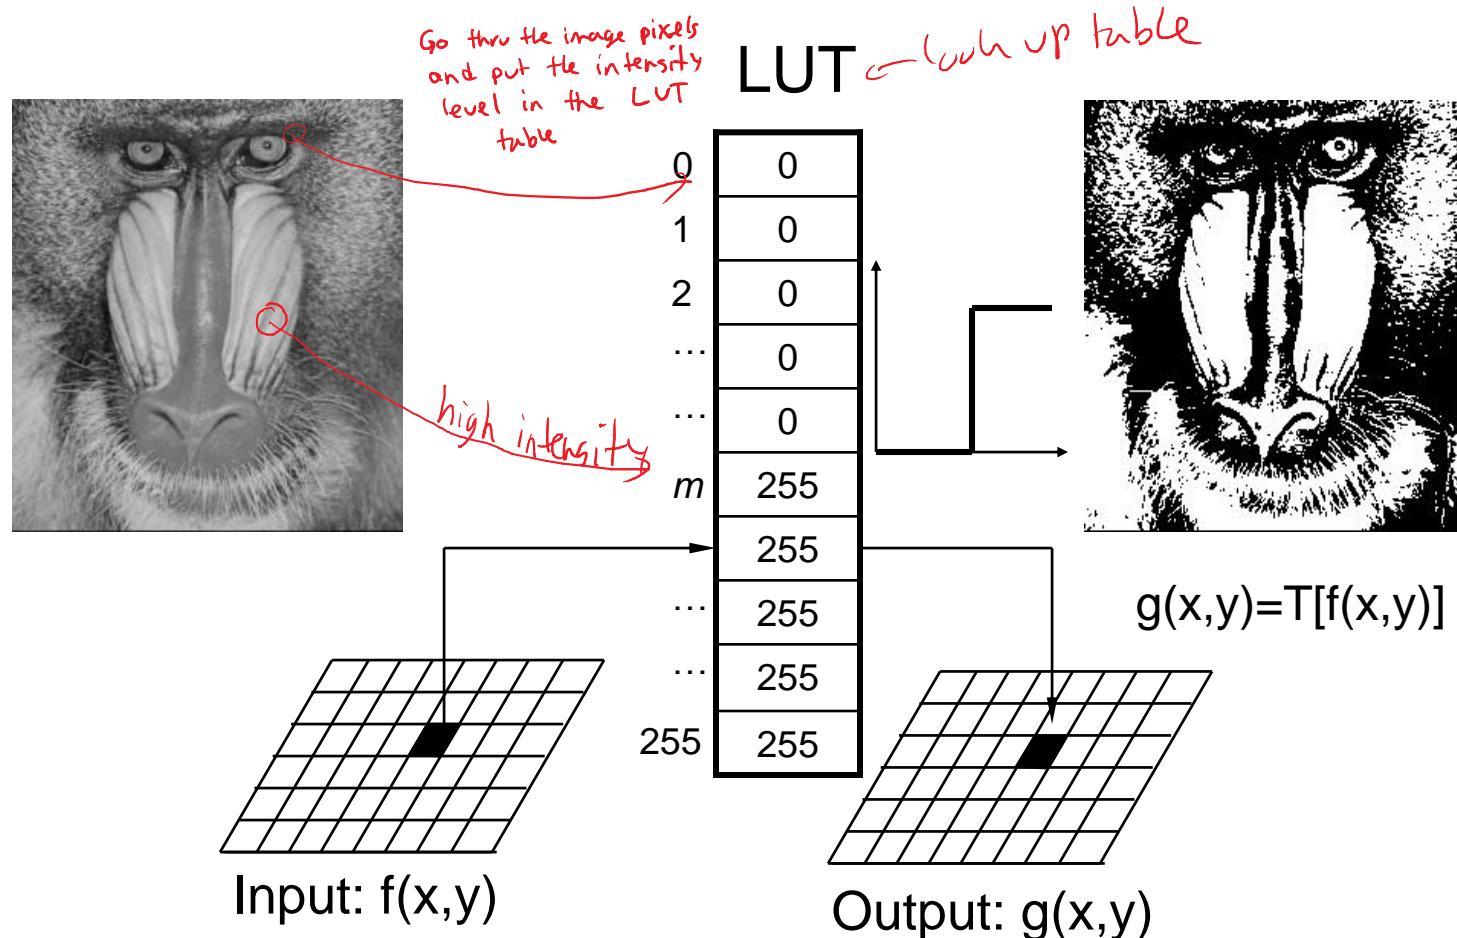
anything below a
certain threshold
is black else
is white

Figure 3.14 An 8-bit grayscale image (left), and the binarized result obtained by thresholding with $\tau = 150$ (right).



Lookup Table: Threshold

- Init LUT with samples taken from thresholding function T



Threshold Program

- Straightforward implementation:

```
// iterate over all pixels
for(i=0; i<total; i++) {
    if(in[i] < thr)      out[i] = BLACK;
    else                  out[i] = WHITE;
}
```

Dont do this way

*If pixel < threshold → black
else → white*

output pixel is :

- Better approach: exploit LUT to avoid **total** comparisons:

```
// init lookup tables
for(i=0; i<thr; i++) lut[i] = BLACK;
for(; i<MXGRAY; i++) lut[i] = WHITE;
```

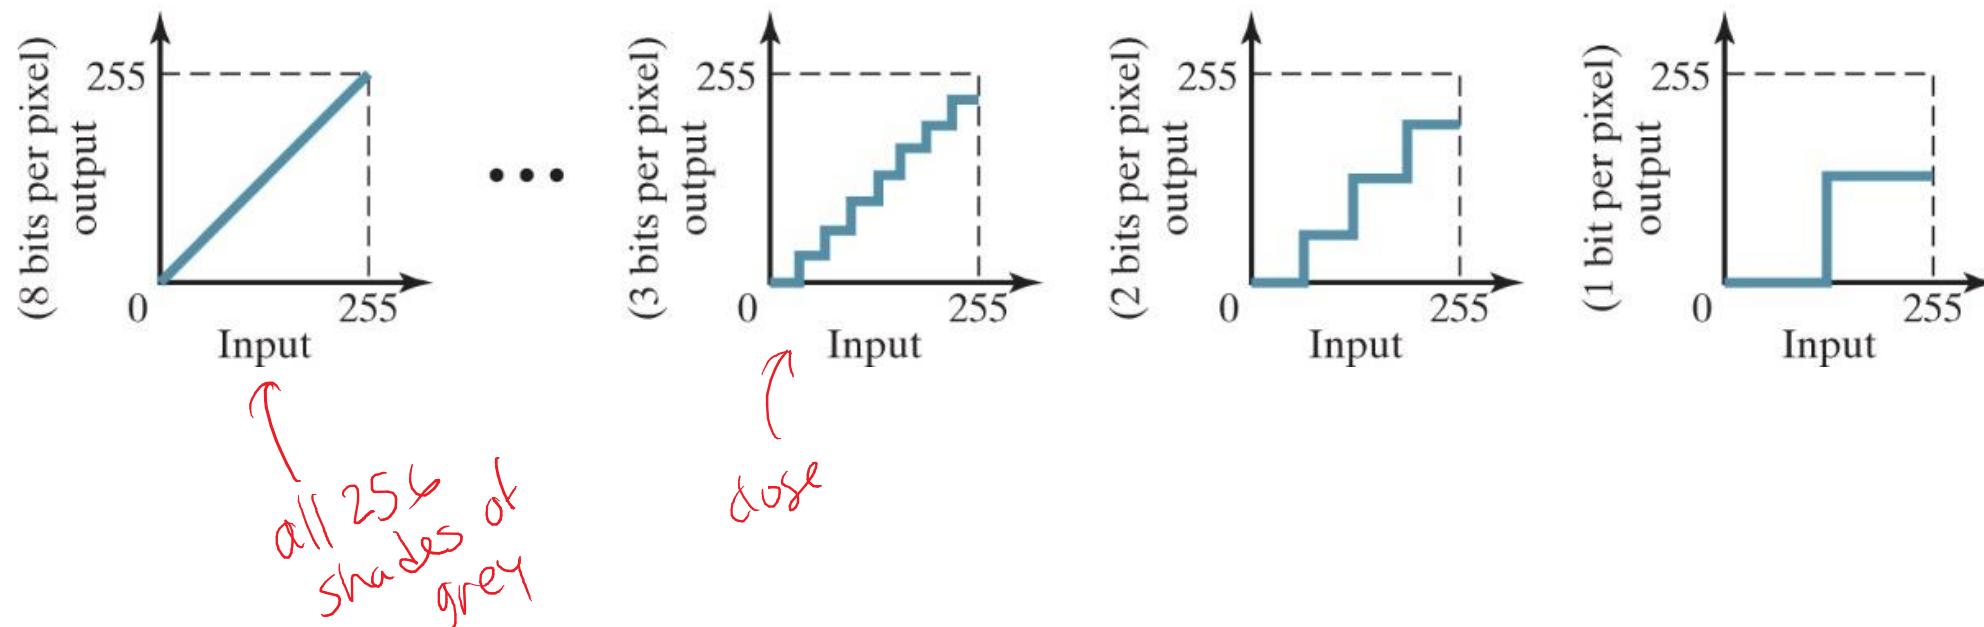
256

assign black or white for look up table

```
// iterate over all pixels
for(i=0; i<total; i++) out[i] = lut[in[i]];
```

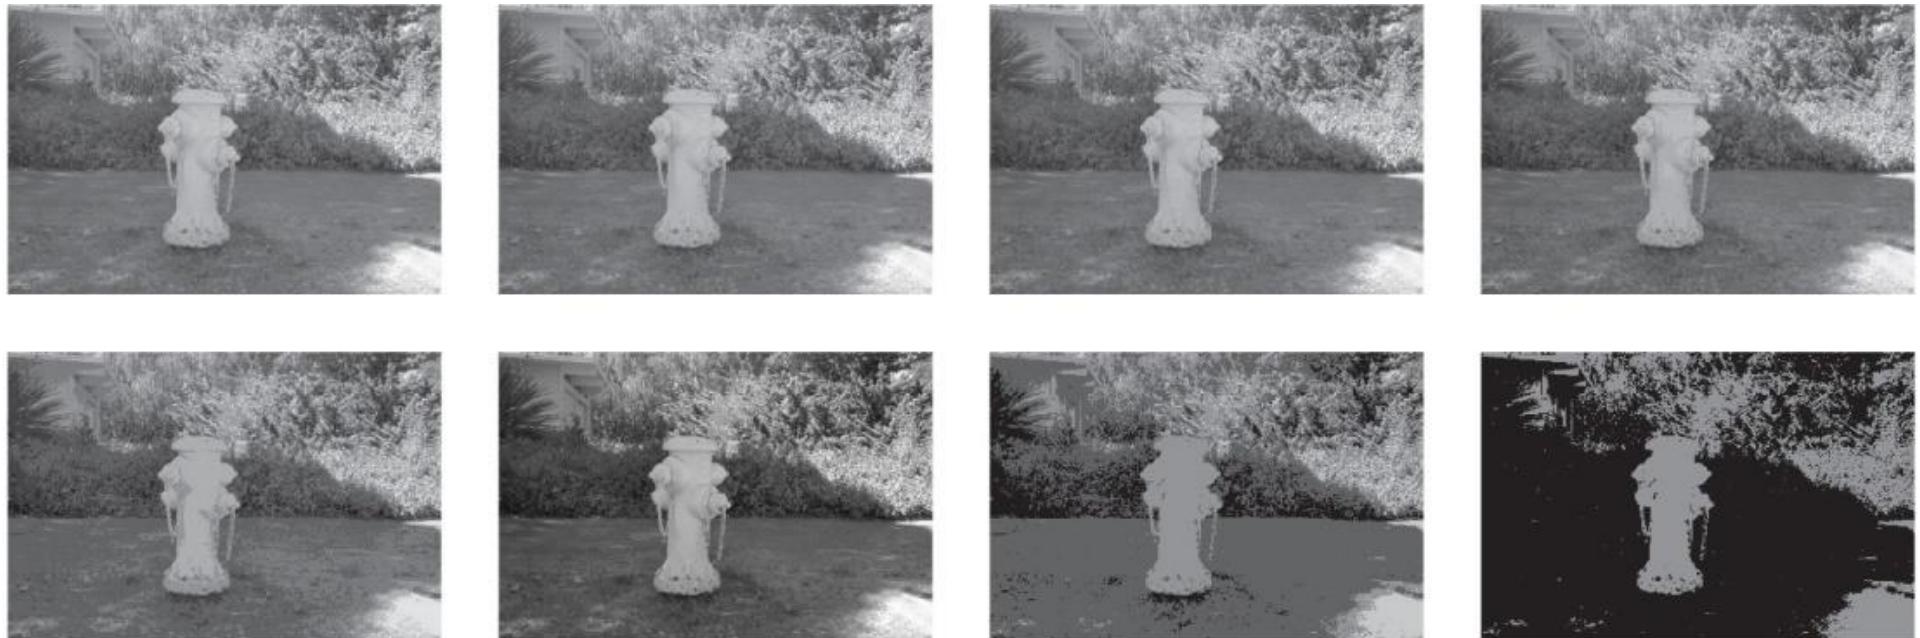
Quantization

Figure 3.16 Quantization discards the lower-order bits via a staircase function, with the number of stairs determined by the number of bits retained. From right to left: Only 1 bit is retained, so the gray levels in the dark half (less than 128) map to 0, while the gray levels in the bright half (above 127) map to 128 (binary: 10000000); 2 bits are retained, so gray levels are mapped to either 0, 64, 128, or 192; 3 bits are retained, so all gray levels are mapped to either 0, 32, 64, 96, 128, 160, 192, or 224; all 8 bits are retained (no quantization, the identity function).



Quantization

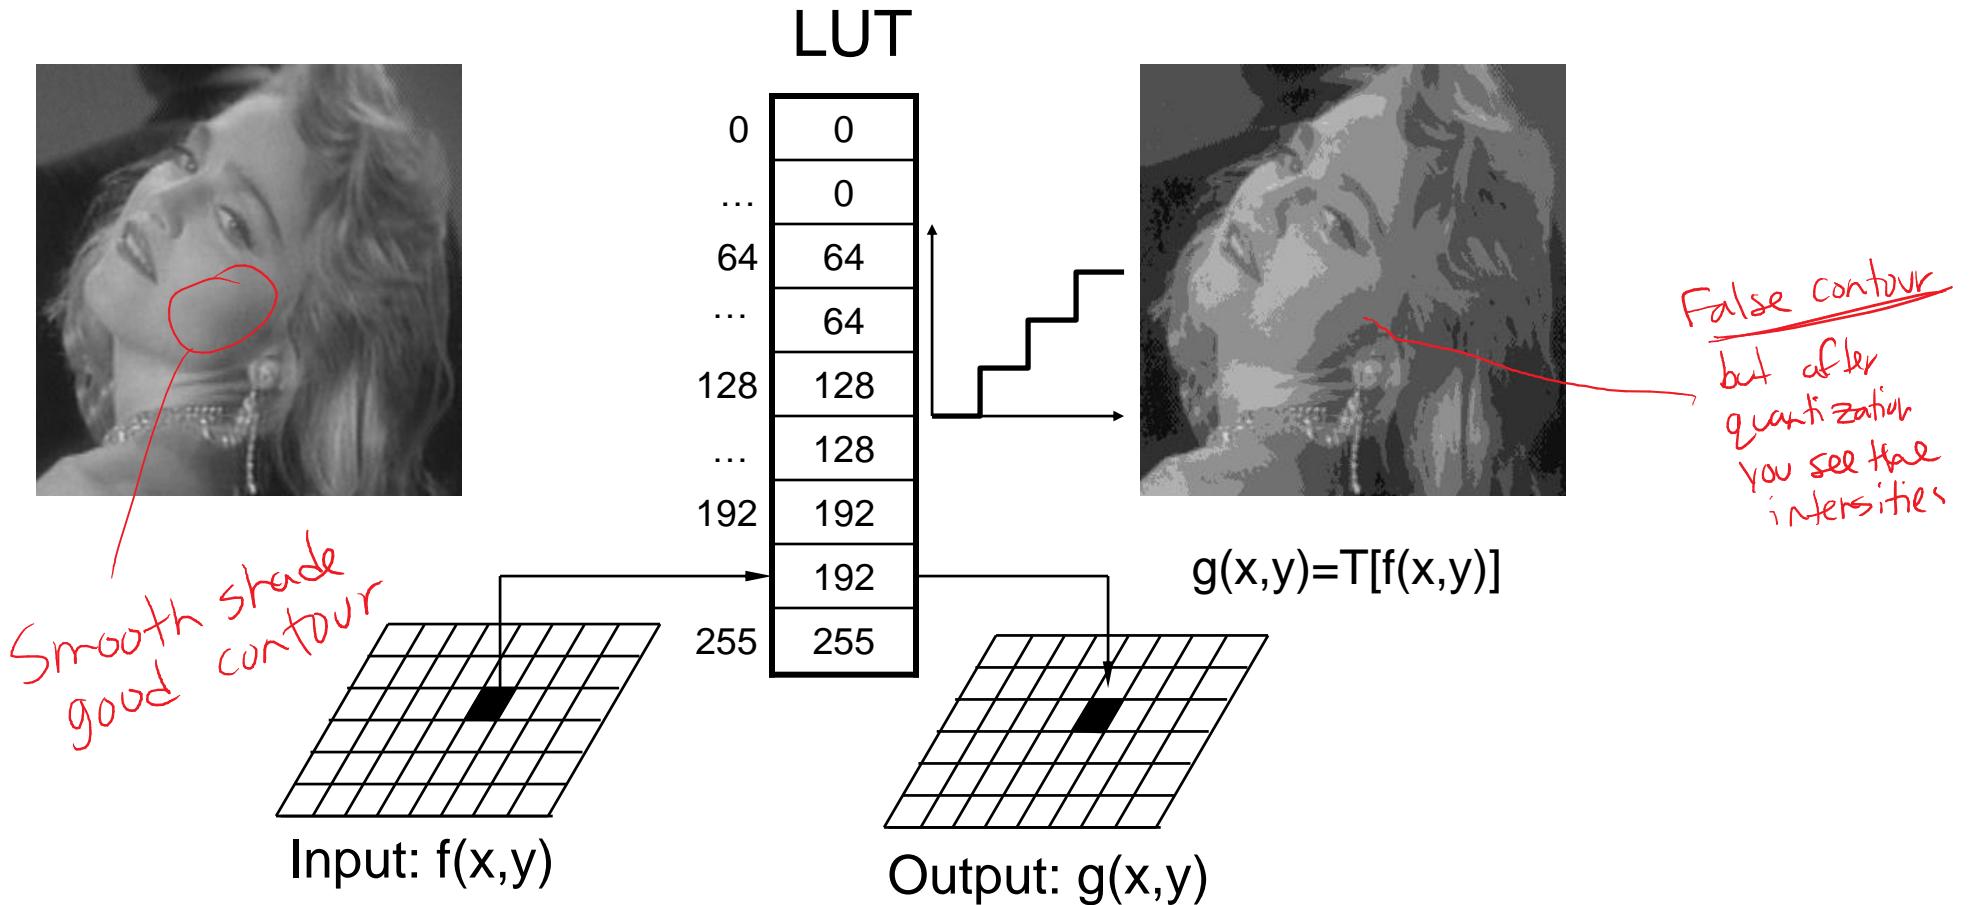
Figure 3.17 From left-to-right and top-to-bottom: An original 8-bit-per-pixel image of a fire hydrant and its quantized versions with 7, 6, 5, 4, 3, 2, and 1 bit per pixel. Note that the image is quite recognizable with as few as 3 bits per pixel.



Stan Birchfield

Lookup Table: Quantization

- Init LUT with samples taken from quantization function T



Quantization Program

- Straightforward implementation:

```
// iterate over all pixels  
scale = MXGRAY / levels; 256/8lut = 32  
for(i=0; i<total; i++)  
    out[i] = scale * (int) (in[i]/scale); all vals 0 to 31 = 0  
                                         2^2 = 4 = 1  
                                         64 - 95 = 2 . . .  
                                         not decimals
```

AVOID its
expensive
for each
pixel

- Better approach: exploit LUT to avoid **total** mults/divisions:

```
// init lookup tables  
scale = MXGRAY / levels;  
for(i=0; i<MXGRAY; i++)  
    lut[i] = scale * (int) (i/scale);
```

done only 256 times

```
// iterate over all pixels  
for(i=0; i<total; i++) out[i] = lut[in[i]];
```

~~Done total lut file.~~

Quantization Artifacts

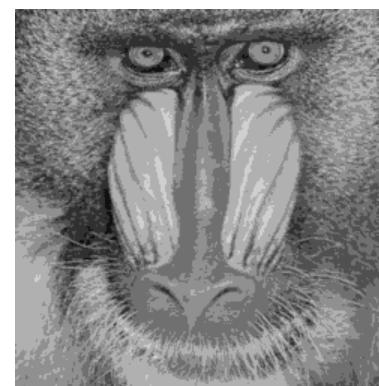
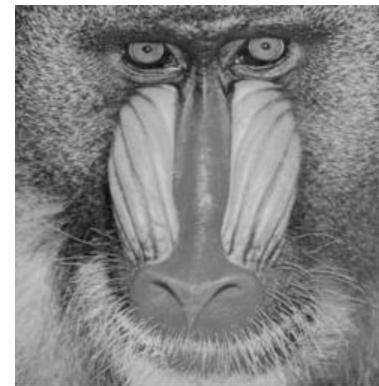
- False contours associated with quantization are most noticeable in smooth areas
- These artifacts are obscured in highly textured regions



Original image



Quantized to 8 levels



Fur is highly textured
↓
So when we quantize it
we won't see the
contours as much
as we see in
smooth surfaces
like faces

*add a little noise
to hide false contours
that arises after
quantization*

Dither Signal

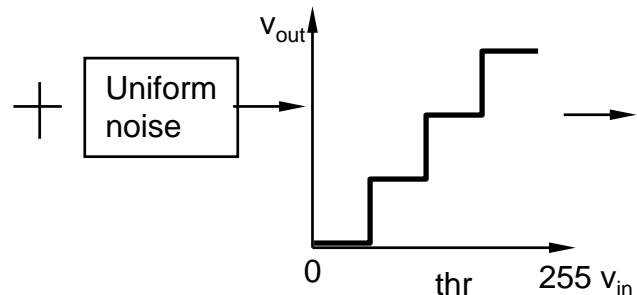
8 Q levels
 $256/8 = 32$

- Reduce quantization error by adding uniformly distributed white noise (dither signal) to the input image prior to quantization.
- Dither hides objectional artifacts.
- To each pixel of the image, add a random number in the range $[-m, m]$, where m is MXGRAY/quantization-levels.

*add an integer to the input
pixel before quantization
doesnt affect look up table*



8 bpp (256 levels)

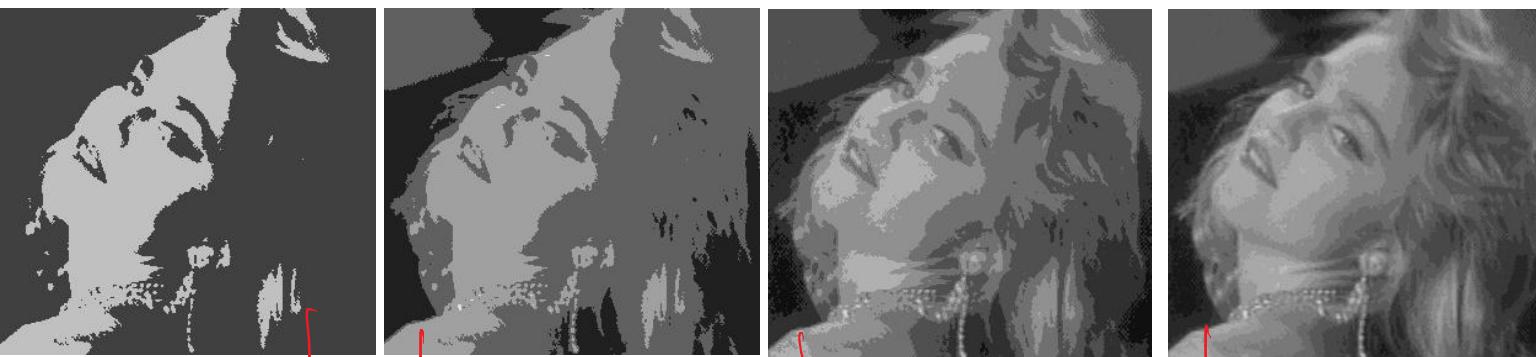


3 bpp (8 levels)

Comparison

1 bit per pixel

Quantization



Dither/
Quantization

With
noise
added!

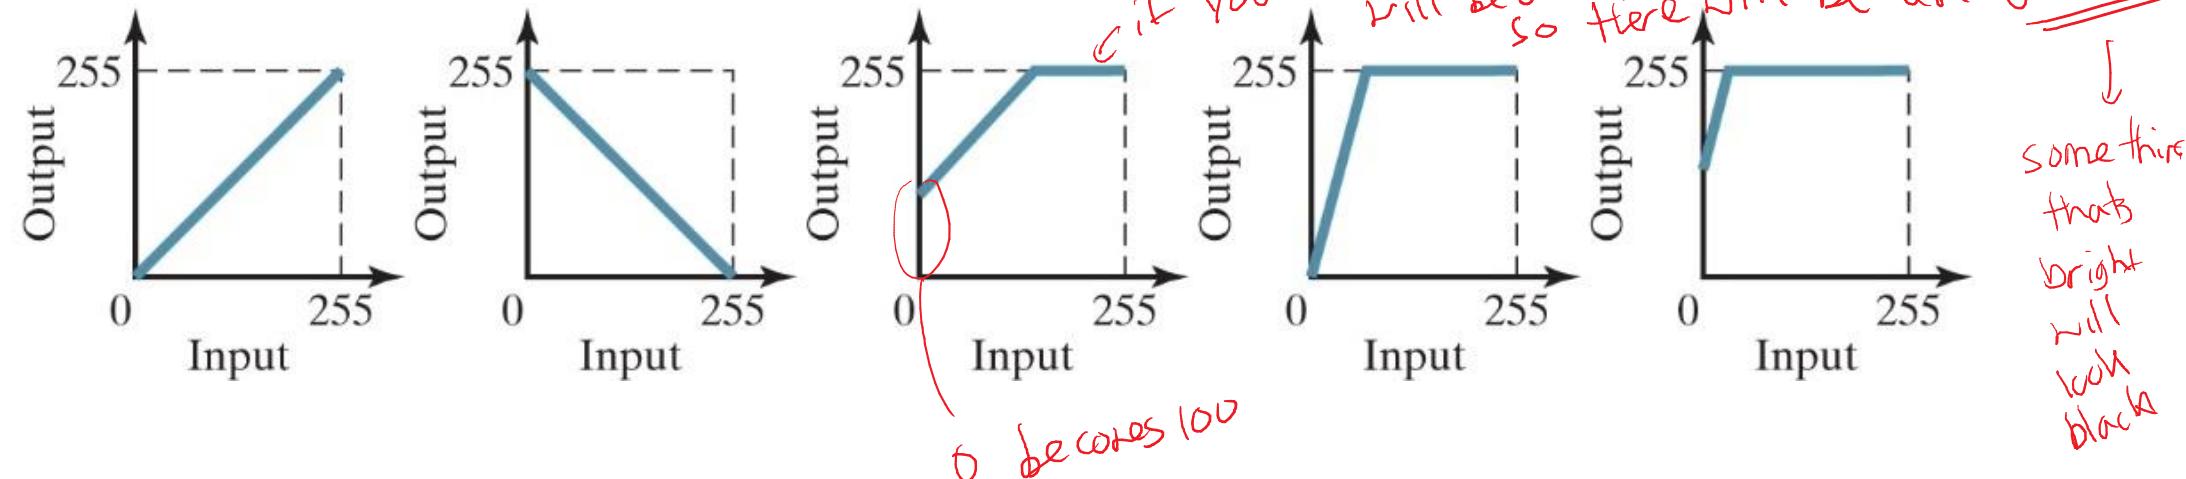


Arithmetic Operations

If you expect a bright pixel
but you see a dark pixel
There is an overflow

- A useful class of graylevel transformations is the set of arithmetic operations, depicted in graphical form in the following figure:

Figure 3.8 Arithmetic graylevel transformations. From left to right: identity, inversion, addition (bias), multiplication (gain), and gain-bias transformation, where saturation arithmetic prevents the output from exceeding the valid range. Note that the slope remains 1 under addition, while the mapping passes through the origin under multiplication.



Saturation Arithmetic

Clamp arithmetic operation to lie in [0, 255] range:

$$[216 \ 171 \ 134 \ 97 \ 52] + 100 = [255 \ 255 \ 234 \ 197 \ 152]$$

$$[216 \ 171 \ 134 \ 97 \ 52] - 100 = [116 \ 71 \ 34 \ 0 \ 0]$$

← Just limit to 0 to 255

ALGORITHM 3.6 Apply gain and bias to a grayscale image, using saturation arithmetic

APPLYGAINANDBIAS(I, b, c)

Input: grayscale image I , constants b, c

Output: grayscale image I' with increased brightness and contrast

```
1 for (x, y) ∈ I do
2      $I'(x, y) \leftarrow \text{MIN}(\text{MAX}(\text{ROUND}(I(x, y) * c + b), 0), 255)$ 
3 return  $I'$ 
```

no decimals
any will be 0
limit to 255
take pixel * c + b
stretch it

Enhancement

- Point operations are used to enhance an image.
- Processed image should be more suitable than the original image for a specific application.
- Suitability is application-dependent.
- A method which is quite useful for enhancing one image may not necessarily be the best approach for enhancing another image.
- Very subjective

Two Enhancement Domains

-
- Spatial Domain: (image plane)
- Techniques are based on direct manipulation of pixels in an image
 - Frequency Domain:
- Techniques are based on modifying the Fourier transform of an image
 - There are some enhancement techniques based on various combinations of methods from these two categories.
- What we were doing by
going thru image → LUT*

Enhanced Images

- For human vision
 - The visual evaluation of image quality is a highly subjective process.
 - It is hard to standardize the definition of a good image.
- For machine perception
 - The evaluation task is easier.
 - A good image is one which gives the best machine recognition results.
- A certain amount of trial and error usually is required before a particular image enhancement approach is selected.

Piecewise-Linear Transformation Functions

- Advantage: The form of piecewise functions can be arbitrarily complex
- Disadvantage: Their specification requires considerably more user input

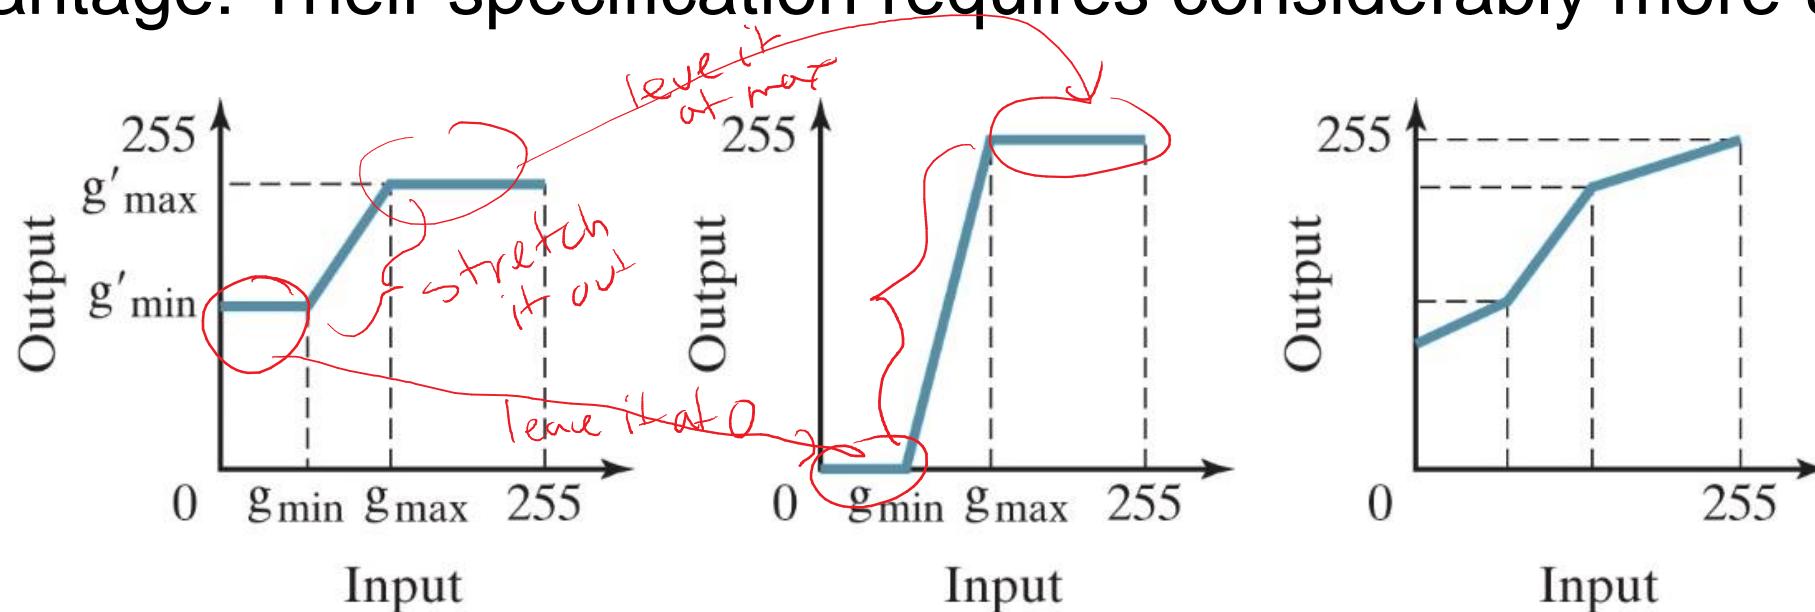
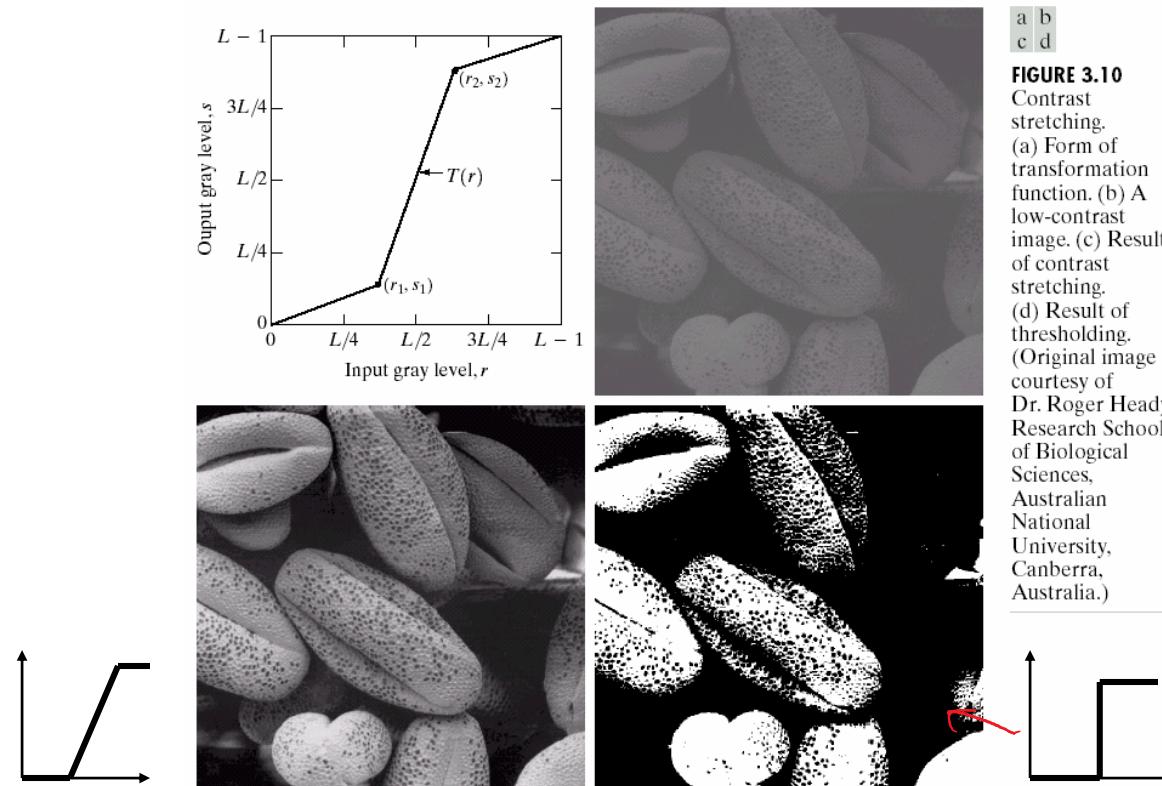


Figure 3.11 Left: Linear contrast stretching maps all the gray levels between g_{\min} and g_{\max} to the range g'_{\min} to g'_{\max} . Middle: If $g'_{\min} = 0$ and $g'_{\max} = 255$, then the full output range is used. Right: A piecewise linear contrast stretch can model any graylevel transformation with arbitrary precision.

Linear Contrast Stretching

- Low contrast may be due to poor illumination, a lack of dynamic range in the imaging sensor, or even a wrong setting of a lens aperture during acquisition.
- Applied contrast stretching: $(r_1, s_1) = (r_{\min}, 0)$ and $(r_2, s_2) = (r_{\max}, L-1)$



a
b
c
d

FIGURE 3.10
Contrast stretching.
(a) Form of transformation function.
(b) A low-contrast image.
(c) Result of contrast stretching.
(d) Result of thresholding.
(Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)

Pick a right graph for good results

Linear Contrast Stretching Transformation

- **Linear contrast stretch:** A transformation that specifies a line segment that maps gray levels between g_{\min} and g_{\max} in the input image to the gray levels g'_{\min} and g'_{\max} in the output image according to a linear function:

$$I'(x, y) = \frac{g'_{\max} - g'_{\min}}{g_{\max} - g_{\min}} (I(x, y) - g_{\min}) + g'_{\min}$$

Maps an image
to gray levels
from g_{\min} to g_{\max}
black \rightarrow white

take out g_{\min}
from image

take input image

Shift
image
to the left

g_{\min} will be mapped
to 0

g_{\min} = map to black
 g_{\max} = map to white

Linear Contrast Stretching Examples

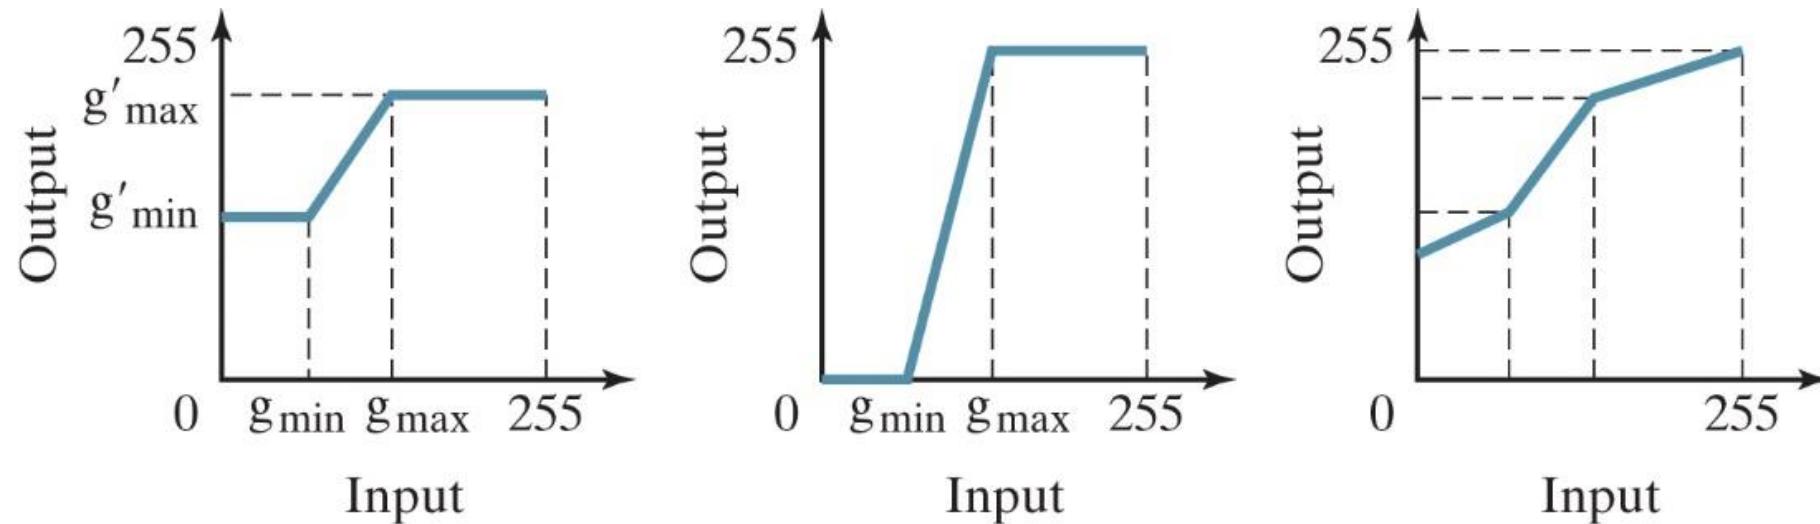


Figure 3.11 Left: Linear contrast stretching maps all the gray levels between g_{\min} and g_{\max} to the range g'_{\min} to g'_{\max} . Middle: If $g'_{\min} = 0$ and $g'_{\max} = 255$, then the full output range is used. Right: A piecewise linear contrast stretch can model any graylevel transformation with arbitrary precision.

Analytic Transformations

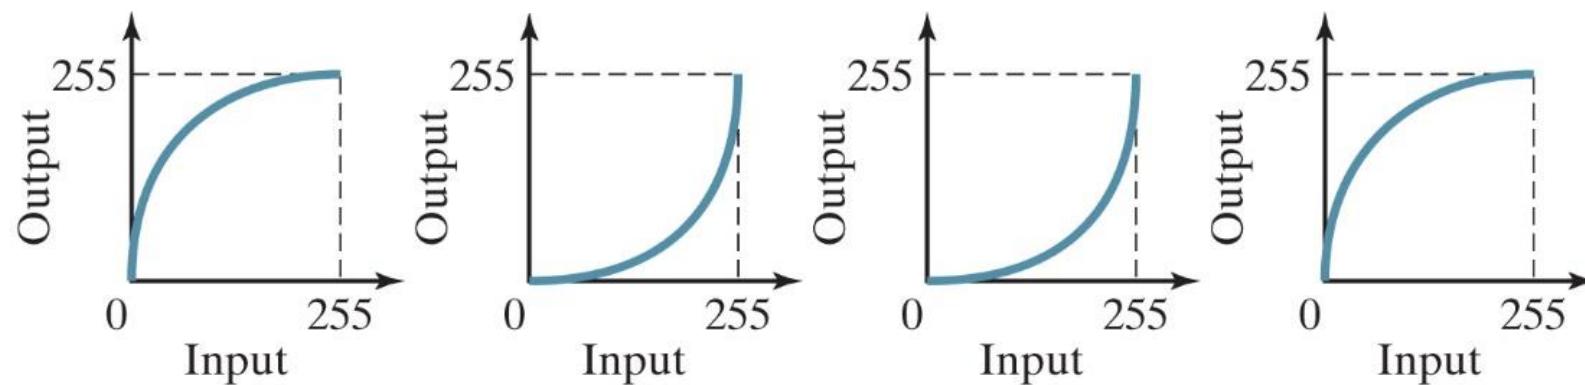
- Graylevel transformations can be specified using analytic functions such as the logarithm, exponential, or power functions:

$$I'(x, y) = \log(I(x, y))$$

$$I'(x, y) = \exp(I(x, y))$$

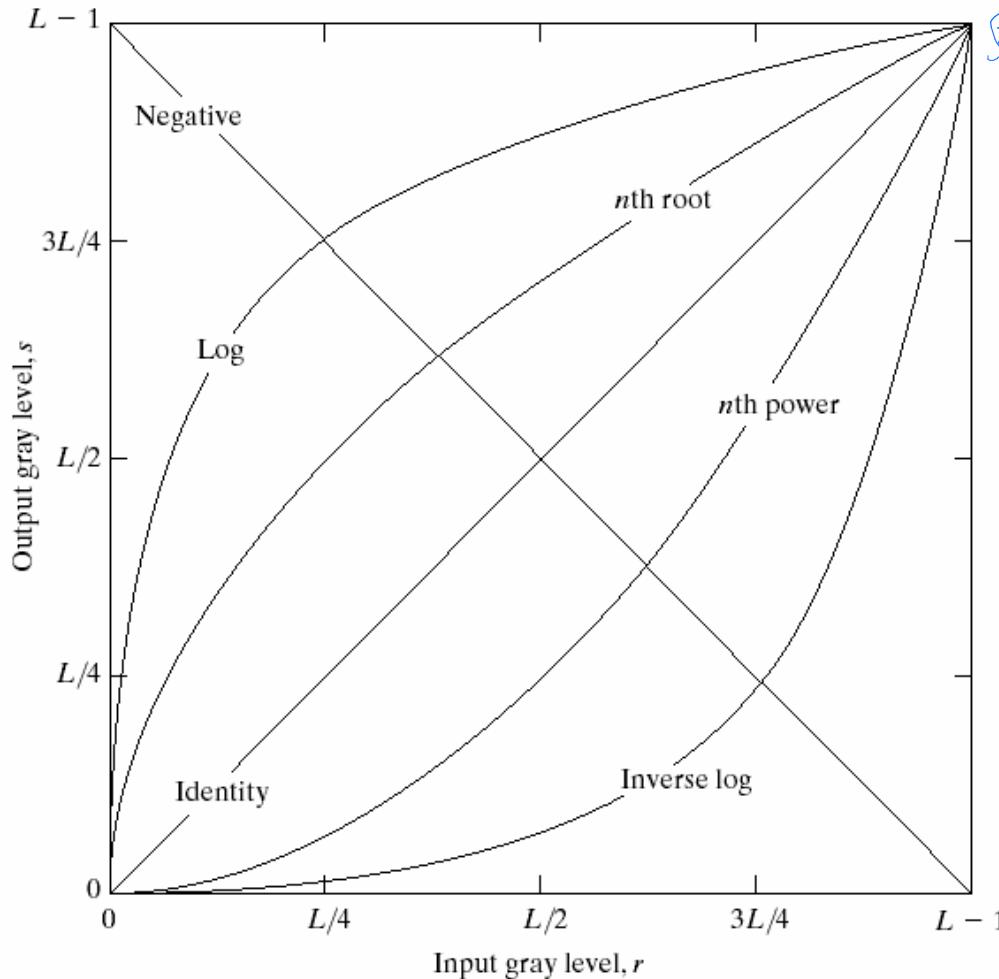
$$I'(x, y) = (I(x, y))^{\gamma}$$

Figure 3.13 Analytic graylevel mapping. From left to right: logarithm, exponential, gamma expansion ($\gamma = 2$), and gamma compression ($\gamma = 0.5$). All transformations are monotonically nondecreasing.



Analytic Transformation Examples

FIGURE 3.3 Some basic gray-level transformation functions used for image enhancement.



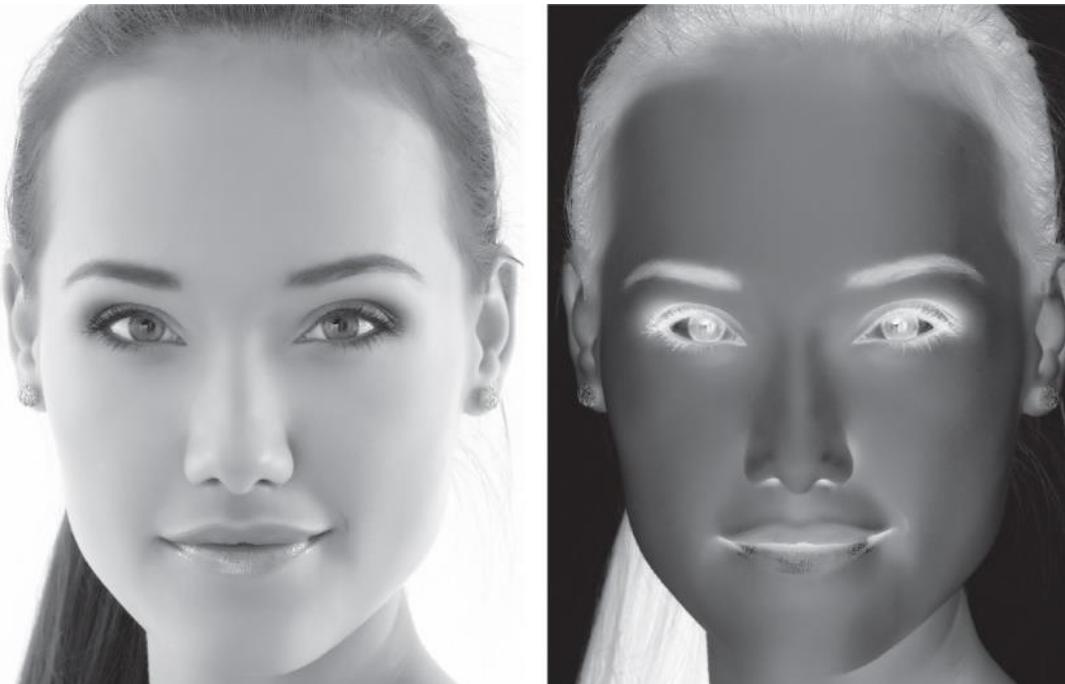
- **Linear function**
 - Negative and identity transformations
- **Logarithmic function**
 - Log and inverse-log transformations
- **Power-law function**
 - n^{th} power and n^{th} root transformations

Normalized ranges from 0 to 1
Can see that because $(L-1)$ raised to
anythis is $(L-1)$
(float not int)
bec remember
 $0.1 \rightarrow 0$ if int

Image Negatives

- Negative transformation : $s = (L-1) - r$
- Reverses the intensity levels of an image.
- Suitable for enhancing white or gray detail embedded in dark regions of an image, especially when black area is large.

Figure 3.9 An 8-bit grayscale image (left), and the inverted image obtained by subtracting each pixel from 255 (right).



$$\log r \Rightarrow 10^? = r$$

Log Transformations

- Log transformation : $s = c \log(1+r)$ ← Why log it?
c is constant and $r \geq 0$
why add +1 to input?
 - Log curve maps a narrow range of low graylevels in input into a wider range of output levels.
 - Expands range of dark image pixels while shrinking bright range.
 - Inverse log expands range of bright image pixels while shrinking dark range.
- $\log(0) = \text{undefined}$ or $\rightarrow \infty$
 $\log(\infty) \approx 0$
so add a 1
this way we won't have an undefined log val.
if $r=0$ $\log(1)=0$

T
useful when we try to
display very large vals

Example of Logarithm Image

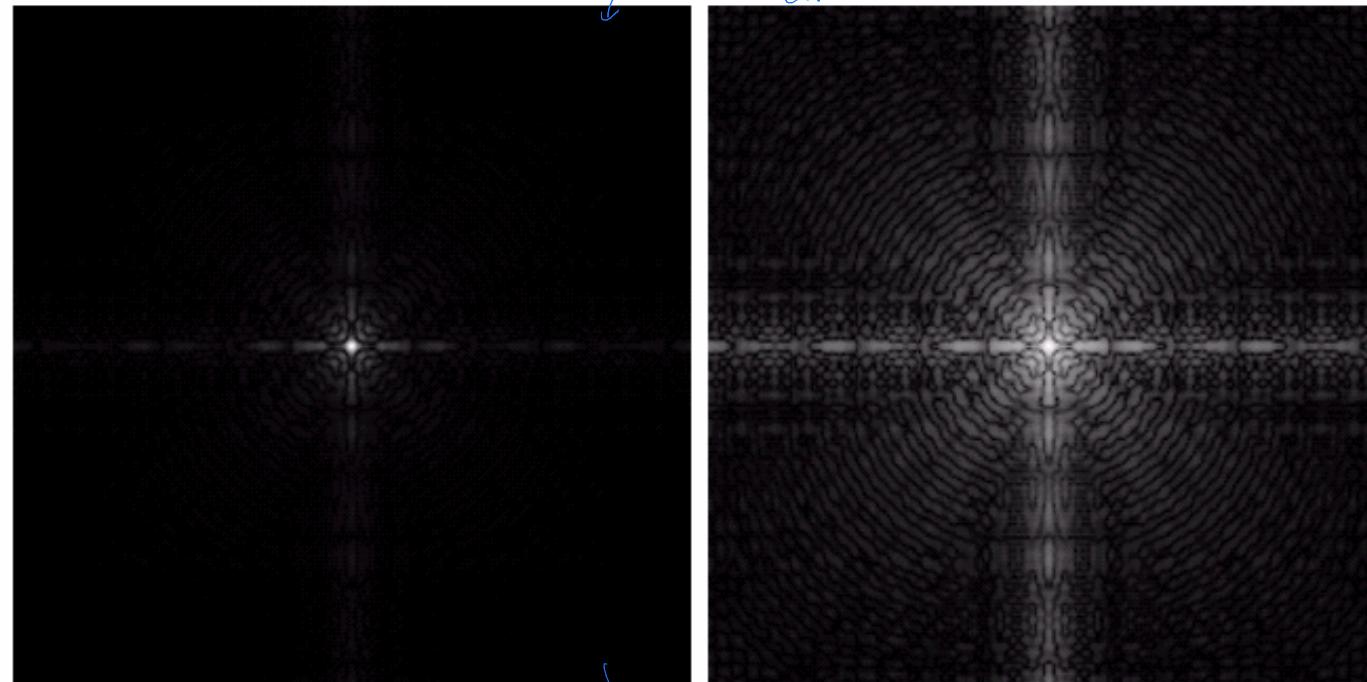
- Fourier spectrum image can have intensity range from 0 to 10^6 or higher.
- Log transform lets us see the detail dominated by large intensity peak.
- Must now display [0,6] range instead of [0, 10^6] range.
- Rescale [0,6] to the [0,255] range.

large pic
to be able to see it, it pushes everything
down to fit and you can't
see it

how can you display a pic
that bigger than the sheet?

a b

FIGURE 3.5
(a) Fourier
spectrum.
(b) Result of
applying the log
transformation
given in
Eq. (3.2-2) with
 $c = 1$.



Power-Law Transformations

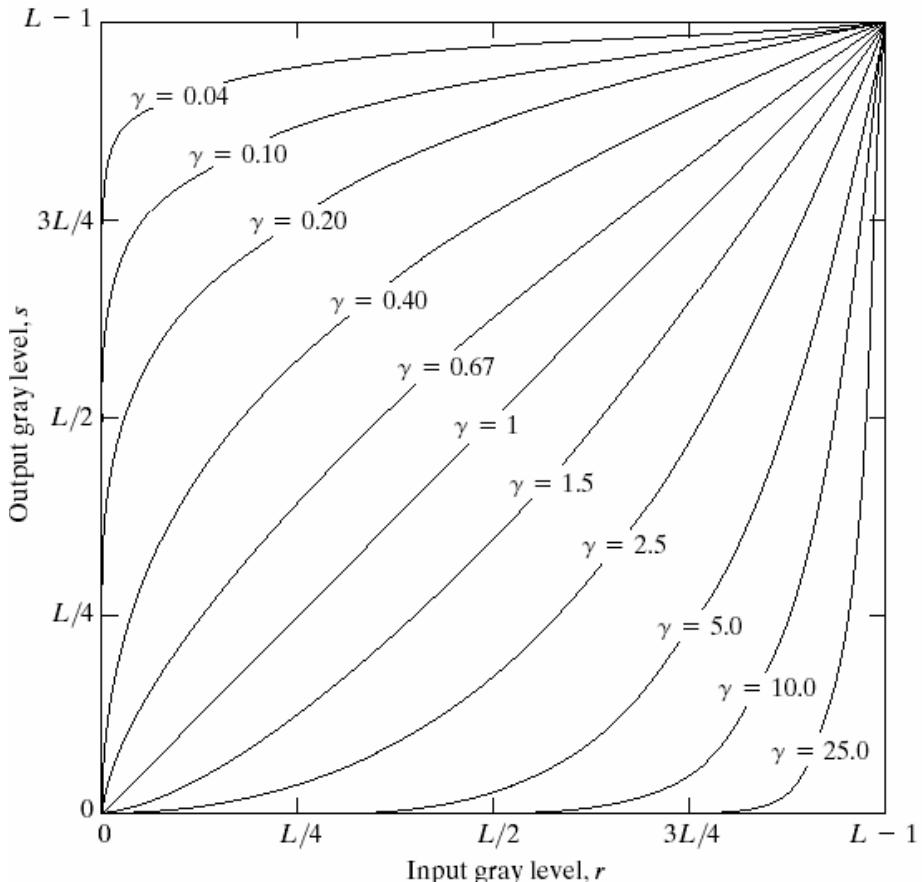


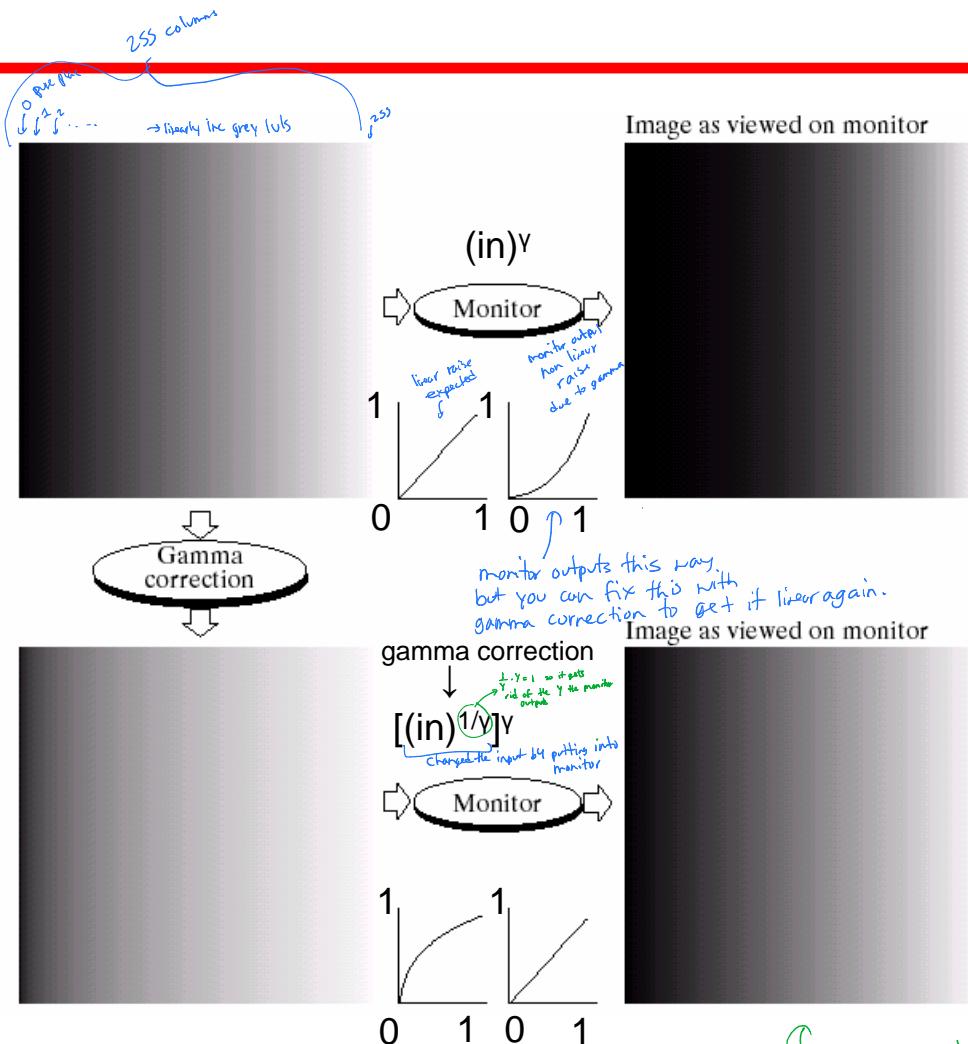
FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases).

- $$s = cr^\gamma$$
- c and γ are positive constants
 - Power-law curves with fractional values of γ map a narrow range of dark input values into a wider range of output values, with the opposite being true for higher values of input levels.
 - $c = \gamma = 1 \Rightarrow$ identity function

Gamma Correction

a
b
c
d

FIGURE 3.7
(a) Linear-wedge gray-scale image.
(b) Response of monitor to linear wedge.
(c) Gamma-corrected wedge.
(d) Output of monitor.



- Cathode ray tube (CRT) devices have an intensity-to-voltage response that is a power function, with γ varying from 1.8 to 2.5
- This darkens the picture.

- Gamma correction is done by preprocessing the image before inputting it to the monitor.

Example: MRI

1st Cure back

2:37

1:57 ← 5 fm

↓
Slide 41

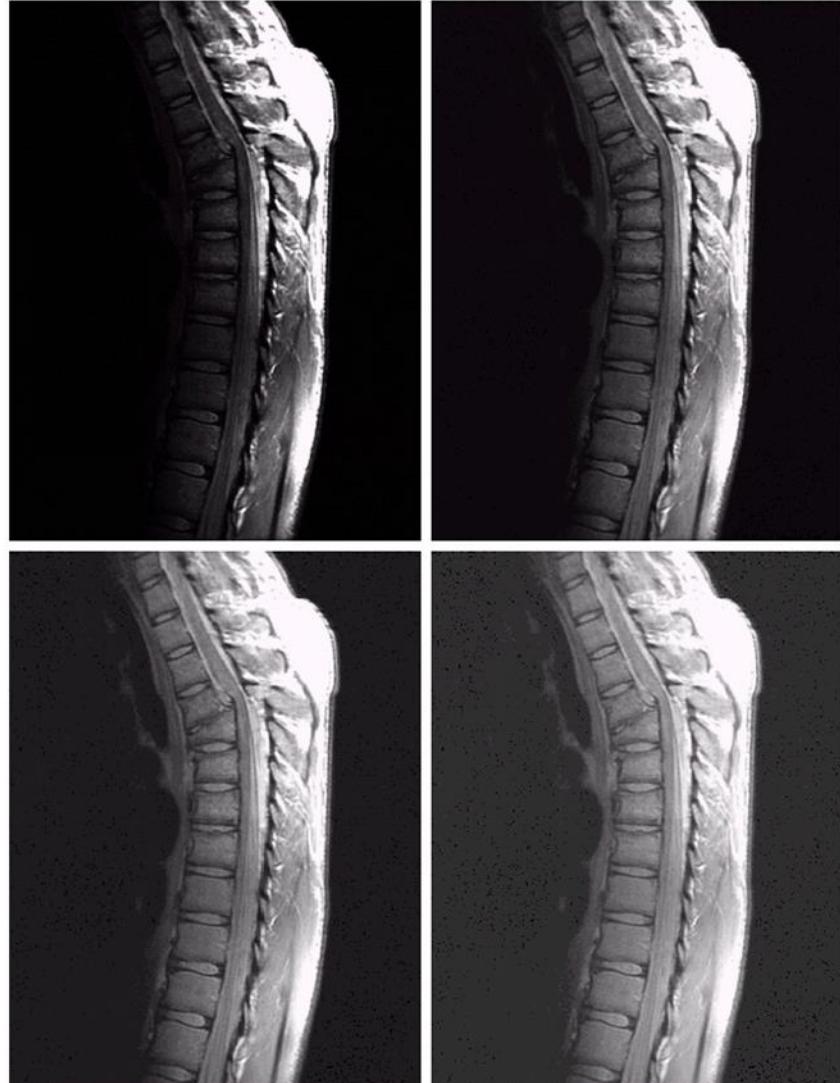


FIGURE 3.8
(a) Magnetic resonance (MR) image of a fractured human spine.
(b)–(d) Results of applying the transformation in Eq. (3.2-3) with $c = 1$ and $\gamma = 0.6, 0.4$, and 0.3 , respectively. (Original image for this example courtesy of Dr. David R. Pickens, Department of Radiology and Radiological Sciences, Vanderbilt University Medical Center.)

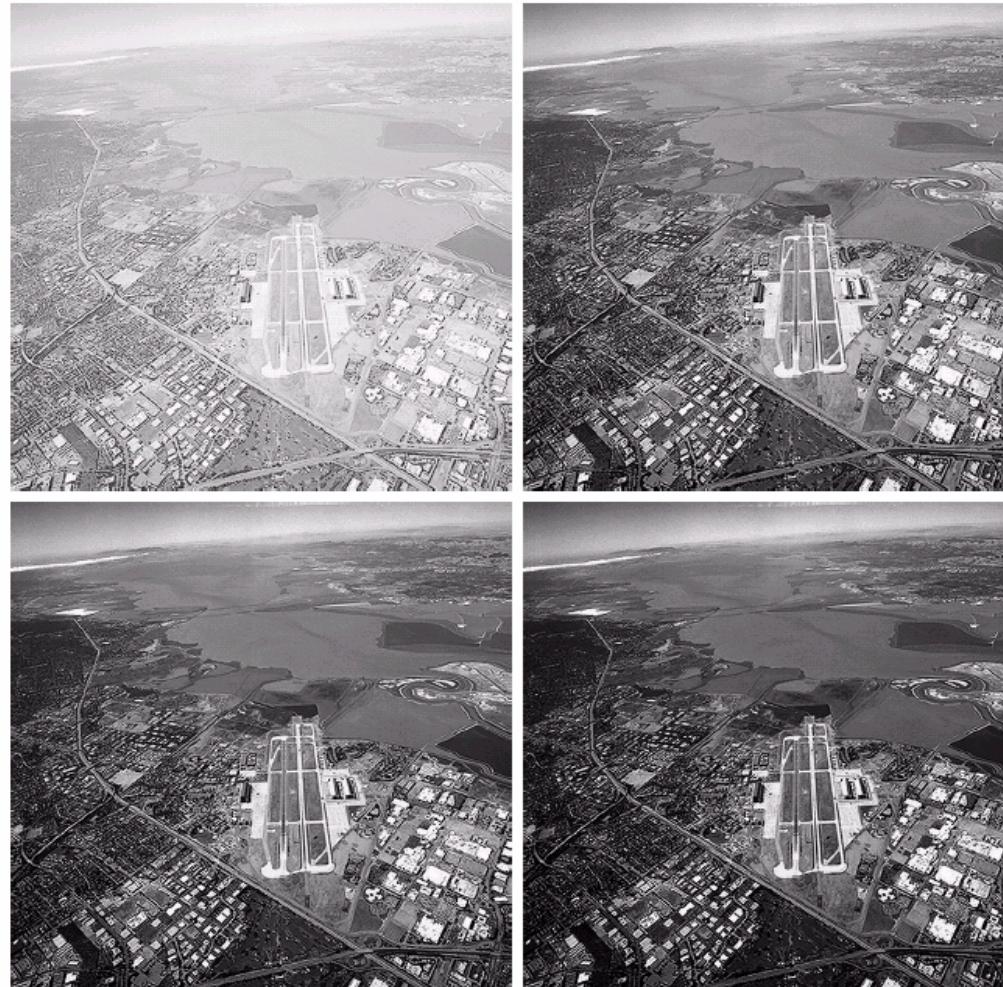
- (a) Dark MRI. Expand graylevel range for contrast manipulation
 $\Rightarrow \gamma < 1$
- (b) $\gamma = 0.6, c=1$
- (c) $\gamma = 0.4$ (best result)
- (d) $\gamma = 0.3$ (limit of acceptability)

When γ is reduced too much, the image begins to reduce contrast to the point where it starts to have a “washed-out” look, especially in the background

Example: Aerial Image

a b
c d

FIGURE 3.9
(a) Aerial image.
(b)–(d) Results of
applying the
transformation in
 $c = 1$ and
 $\gamma = 3.0, 4.0,$ and
 $5.0,$ respectively.
(Original image
for this example
courtesy of
NASA.)



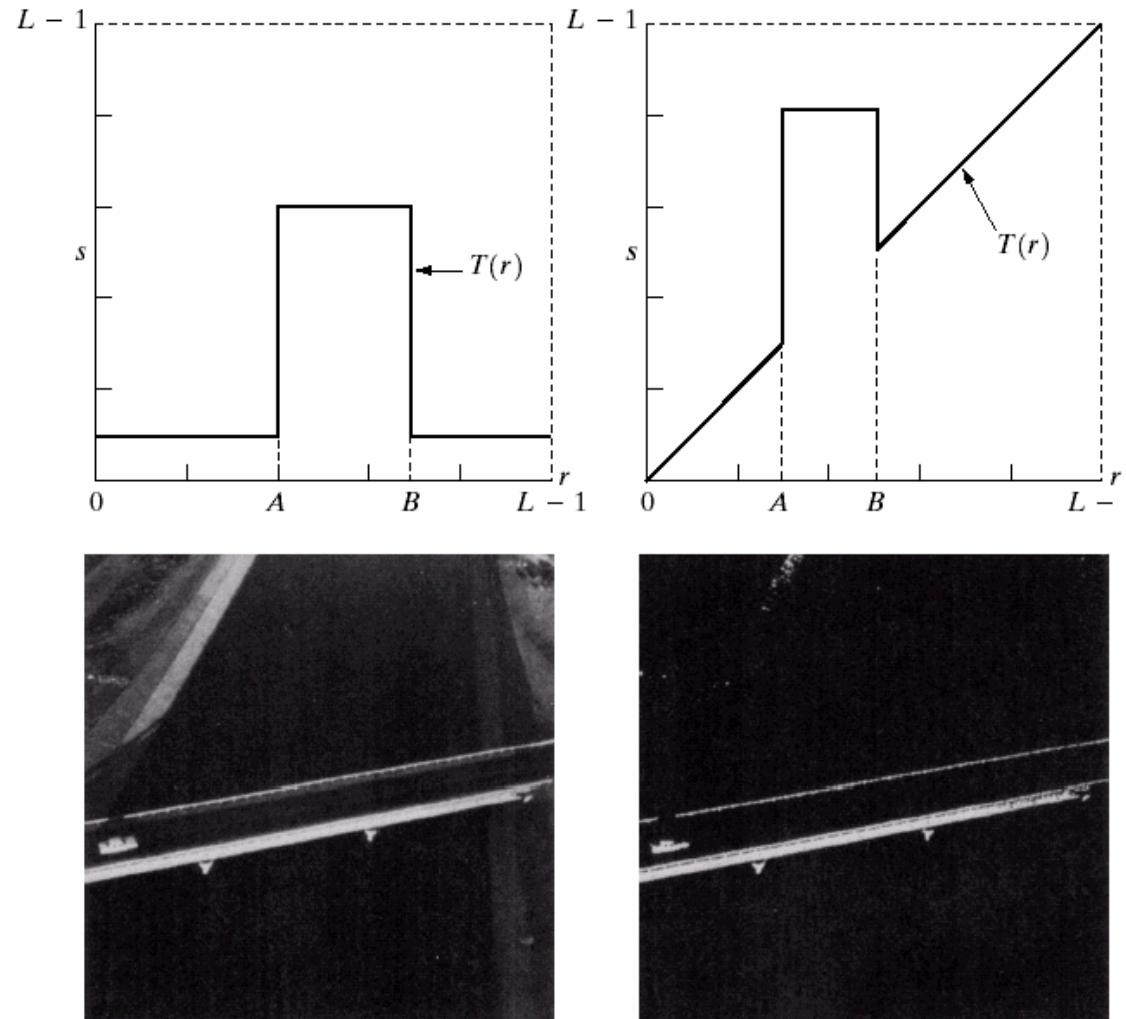
Washed-out
image. Shrink
graylevel range
 $\Rightarrow \gamma > 1$

(b) $\gamma = 3.0$
(suitable)

(c) $\gamma = 4.0$
(suitable)

(d) $\gamma = 5.0$
(High contrast;
the image has
areas that are
too dark; some
detail is lost)

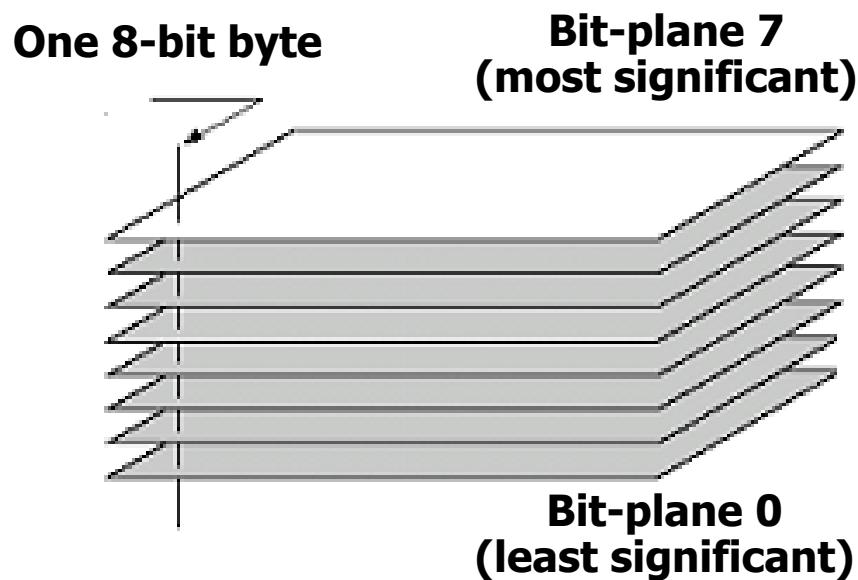
Graylevel Slicing



a b
c d

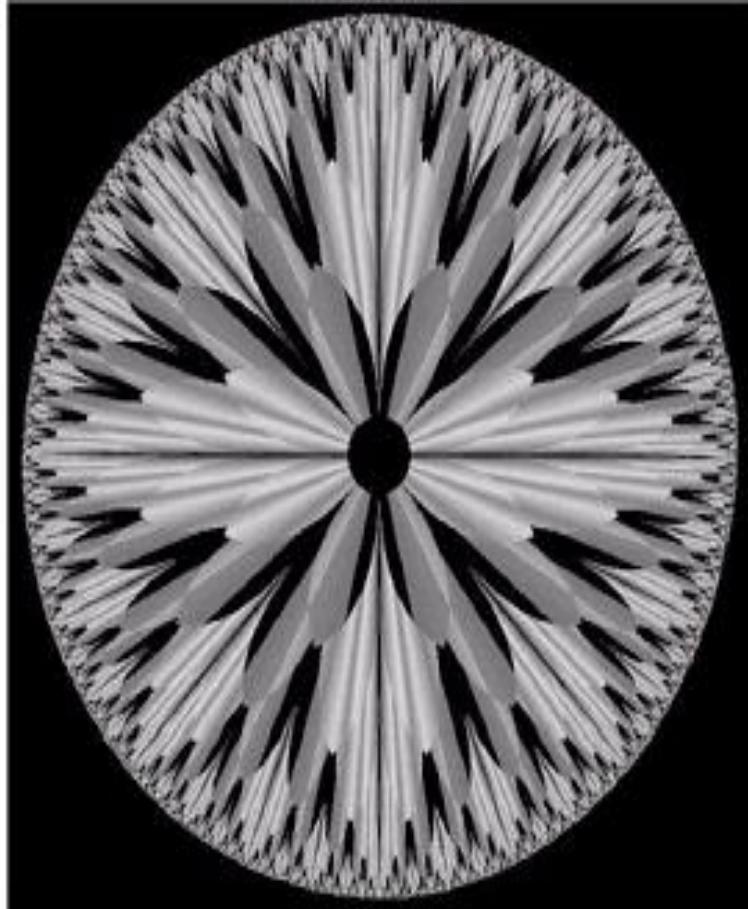
FIGURE 3.11
(a) This transformation highlights range $[A, B]$ of gray levels and reduces all others to a constant level.
(b) This transformation highlights range $[A, B]$ but preserves all other levels.
(c) An image.
(d) Result of using the transformation in (a).

Bit-plane slicing



- Highlighting the contribution made to total image appearance by specific bits
- Suppose each pixel is represented by 8 bits
Higher-order bits contain the majority of the visually significant data
Useful for analyzing the relative importance played by each bit of the image

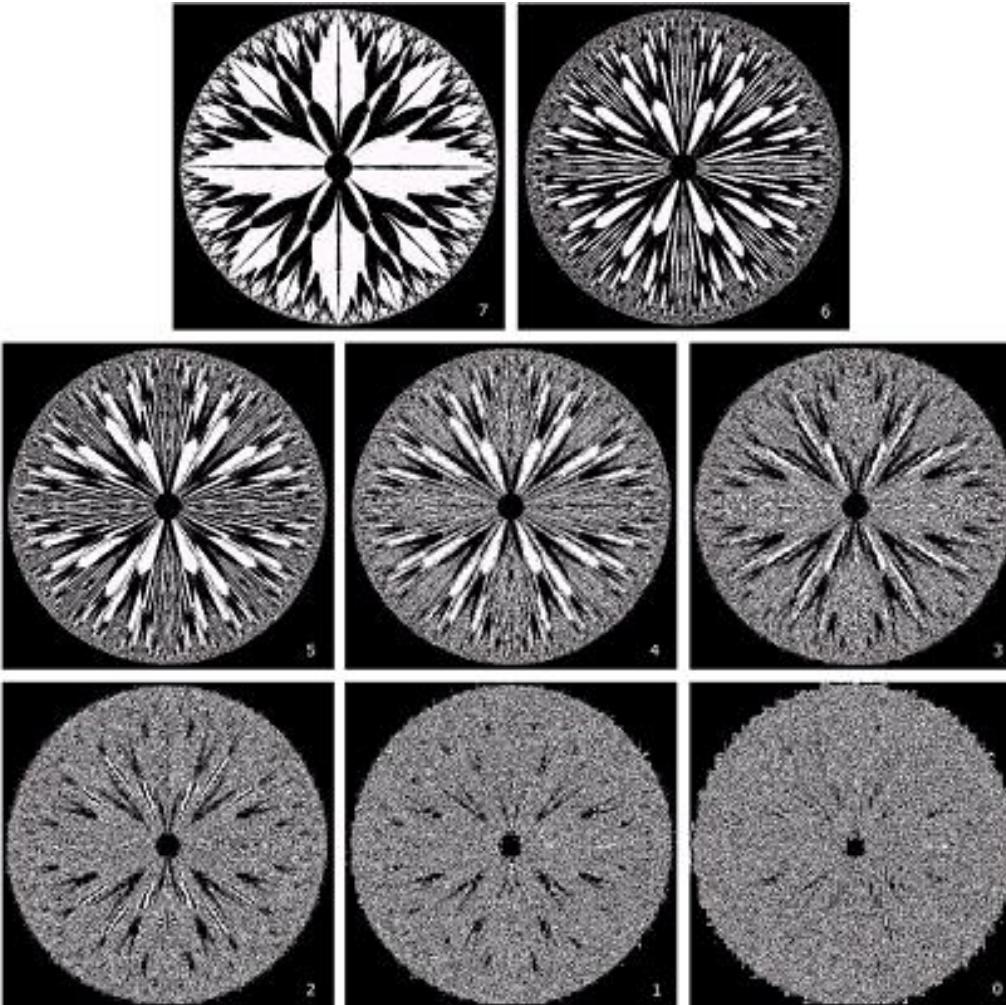
Example



An 8-bit fractal image

- The (binary) image for bit-plane 7 can be obtained by processing the input image with a thresholding graylevel transformation.
 - Map all levels between 0 and 127 to 0
 - Map all levels between 129 and 255 to 255

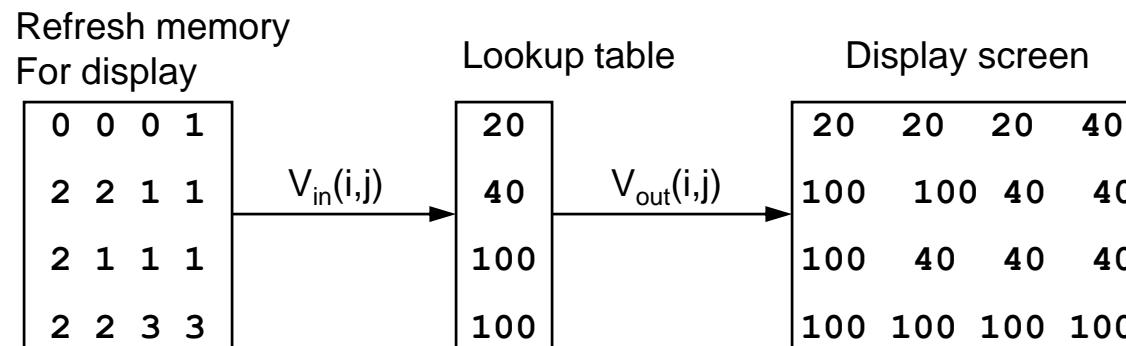
8-Bit Planes



	Bit-plane 7	Bit-plane 6
Bit-plane 5	Bit-plane 4	Bit-plane 3
Bit-plane 2	Bit-plane 1	Bit-plane 0

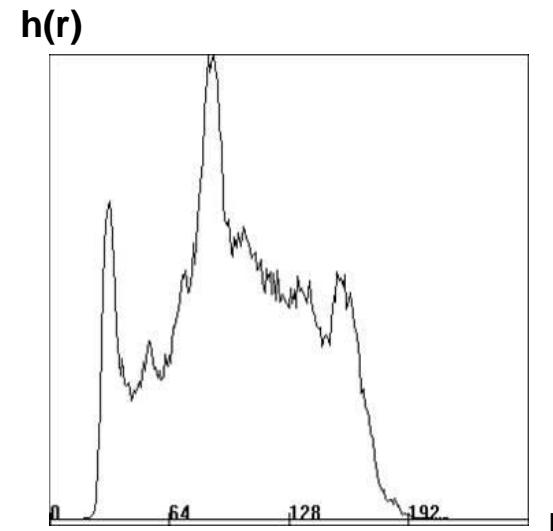
Hardware LUTs

- All point operations can be implemented by LUTs.
- Hardware LUTs operate on the data as it is being displayed.
- It's an efficient means of applying transformations because changing display characteristics only requires loading a new table and not the entire image.
- For a 1024x1024 8-bit image, this translates to 256 entries instead of one million.
- LUTs do not alter the contents of original image (nondestructive).



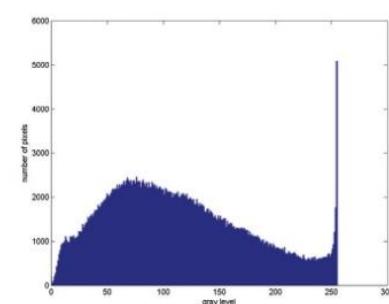
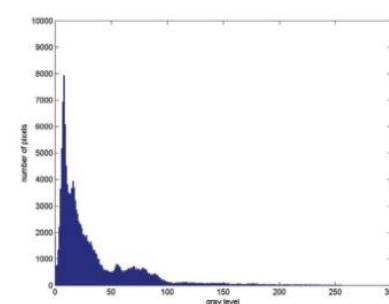
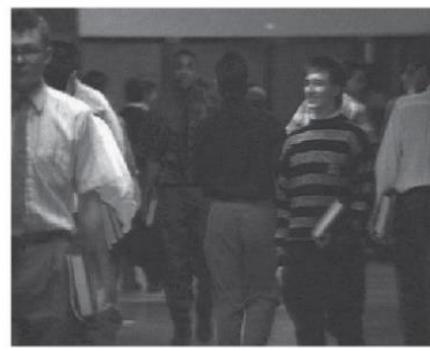
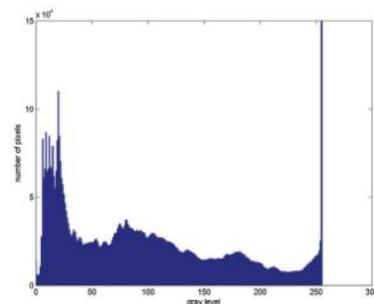
Histogram

- A histogram of a digital image with gray levels in the range $[0, L-1]$ is a discrete function $h(r_k) = n_k$
 - r_k : the k^{th} gray level
 - n_k : the number of pixels in the image having gray level r_k
- The sum of all histogram entries is equal to the total number of pixels in the image.



Histogram Examples

Figure 3.20 Images and their histograms. From left to right: an image with high contrast and many dark or bright pixels, a dark image with low contrast, and another high contrast image with good exposure. Note the spikes at 255 in the first and last images, indicating pixel saturation.



Stan Birchfield, Source: Movie Hoop Dreams, Jessica Birchfield

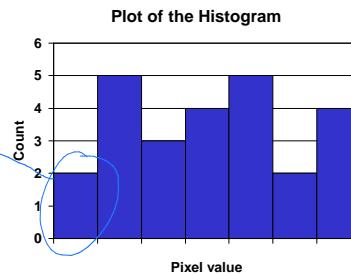
Histogram Evaluation

5x5 image

2	3	4	4	6
1	2	4	5	6
1	1	5	6	6
0	1	3	3	4
0	1	2	3	4

for gray val 0
; found 2

Graylevel	Count
0	2
1	5
2	3
3	4
4	5
5	2
6	4
Total	25



basically increments
these vals

Histogram evaluation:

```
for(i=0; i<MXGRAY; i++) H[i] = 0;  
for(i=0; i<total; i++) H[in[i]]++;
```

input my value
index i+

Normalized Histogram

- Divide each histogram entry at gray level r_k by the total number of pixels in the image, n

$$H(r) = h$$

$$p(r_k) = n_k / n \quad \leftarrow \text{What's the probability there is a gray val at a specific pixel?}$$

- $p(r_k)$ gives an estimate of the probability of occurrence of gray level r_k
- The sum of all components of a normalized histogram is equal to 1.

Sum of non normalized = total # of pixel

but we divided by n

so

$$\text{normalized} = 1 : \cdot \cdot \cdot$$

Pseudocode

h = Length of image ?
 H = histogram

ALGORITHM 3.7 Compute the graylevel histogram of an image

COMPUTEHISTOGRAM(I)

Input: grayscale image I

Output: graylevel histogram h

1 for $\ell \leftarrow 0$ to 255 do

2 $h[\ell] \leftarrow 0$

3 for $(x, y) \in I$ do

4 $h[I(x, y)] \leftarrow +1$ *let it be incremented by 1*

$\triangleright h[I(x, y)] \leftarrow h[I(x, y)] + 1$

5 return h

ALGORITHM 3.8 Compute the normalized graylevel histogram of an image

COMPUTENORMALIZEDHISTOGRAM(I)

Input: grayscale image I

Output: normalized graylevel histogram \bar{h}

1 $h \leftarrow \text{COMPUTEHISTOGRAM}(I)$

2 $n \leftarrow \text{width} * \text{height}$

3 for $\ell \leftarrow 0$ to 255 do

4 $\bar{h}[\ell] \leftarrow h[\ell]/n$ *take each histogram entry and divide by n*

5 return \bar{h}

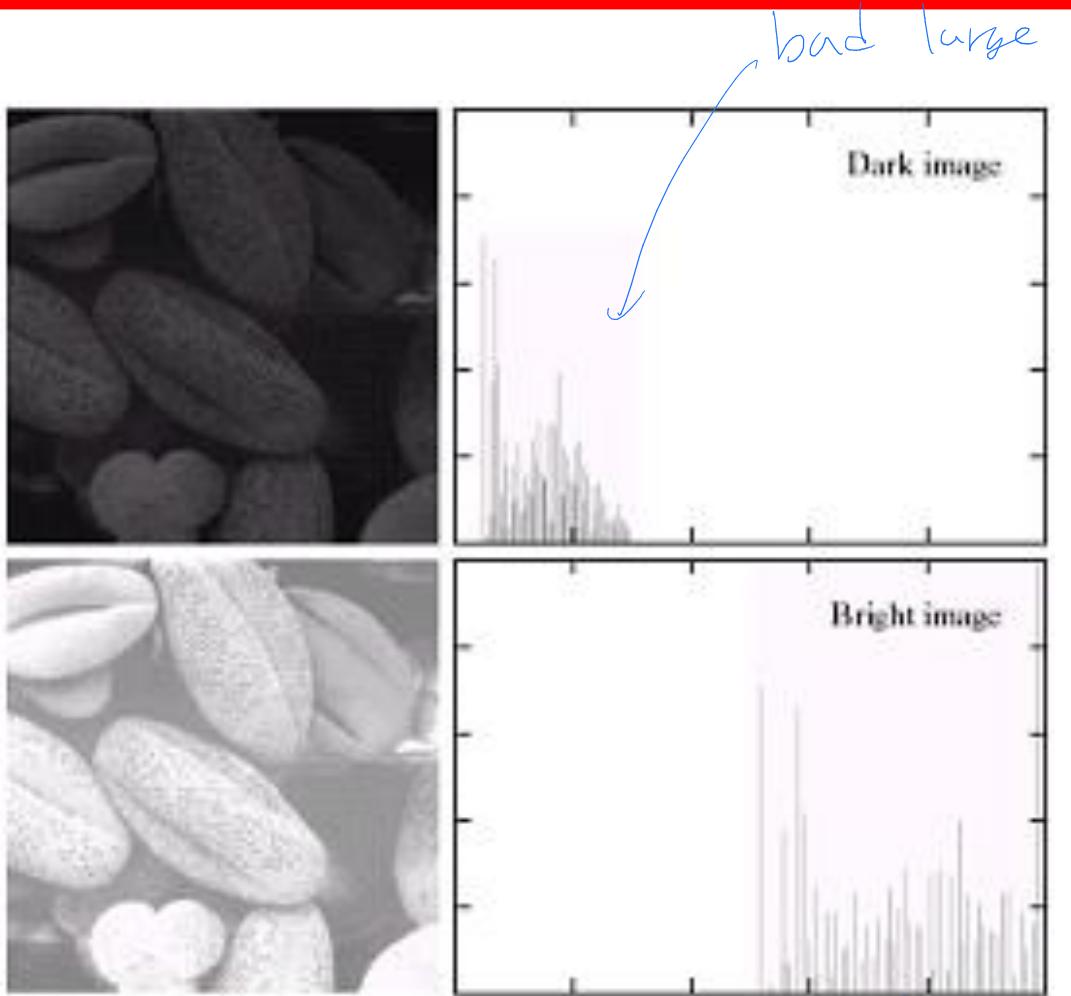
and assign to a new one called \bar{h}

Histogram Processing

- Basic for numerous spatial domain processing techniques.
- Used effectively for image enhancement:
 - Histogram stretching
 - Histogram equalization
 - Histogram matching

Will use this to make histogram image better
- Information inherent in histograms also is useful in image compression and segmentation.

Example: Dark/Bright Images



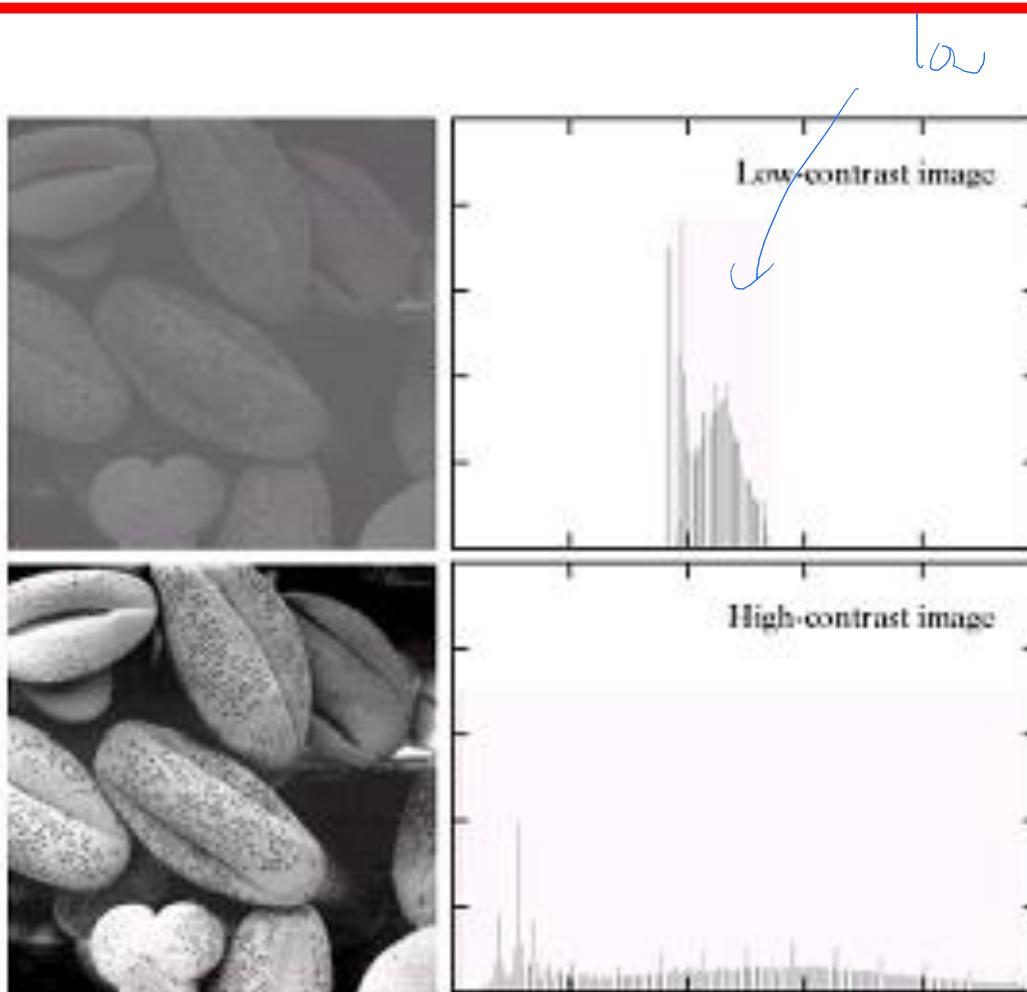
Dark image

Components of histogram are concentrated on the low side of the gray scale.

Bright image

Components of histogram are concentrated on the high side of the gray scale.

Example: Low/High Contrast Images



low contrast image

Low-contrast image

histogram is narrow
and centered toward
the middle of the
gray scale

High-contrast image

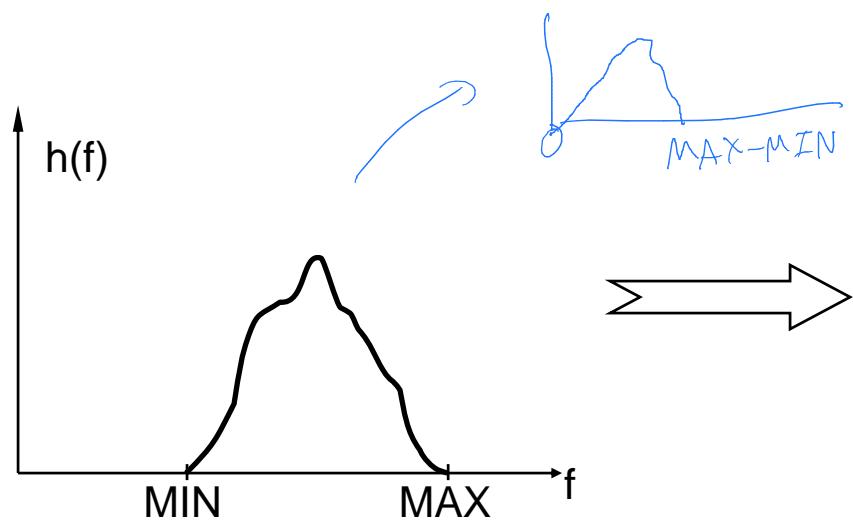
histogram covers broad
range of the gray scale
and the distribution of
pixels is not too far from
uniform, with very few
vertical lines being much
higher than the others

How to get from low contrast to high contrast? Here is one technique.

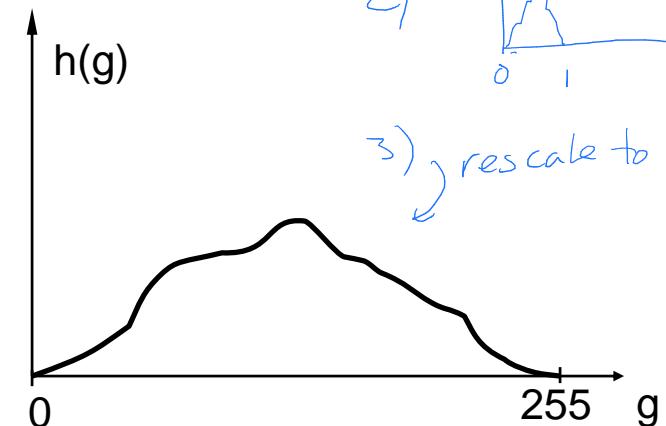


Histogram Stretching

1) slide histogram to 0 (subtract MIN from each pixel)



2) normalize to 0 to 1 range



3) rescale to 0 to 255 range.

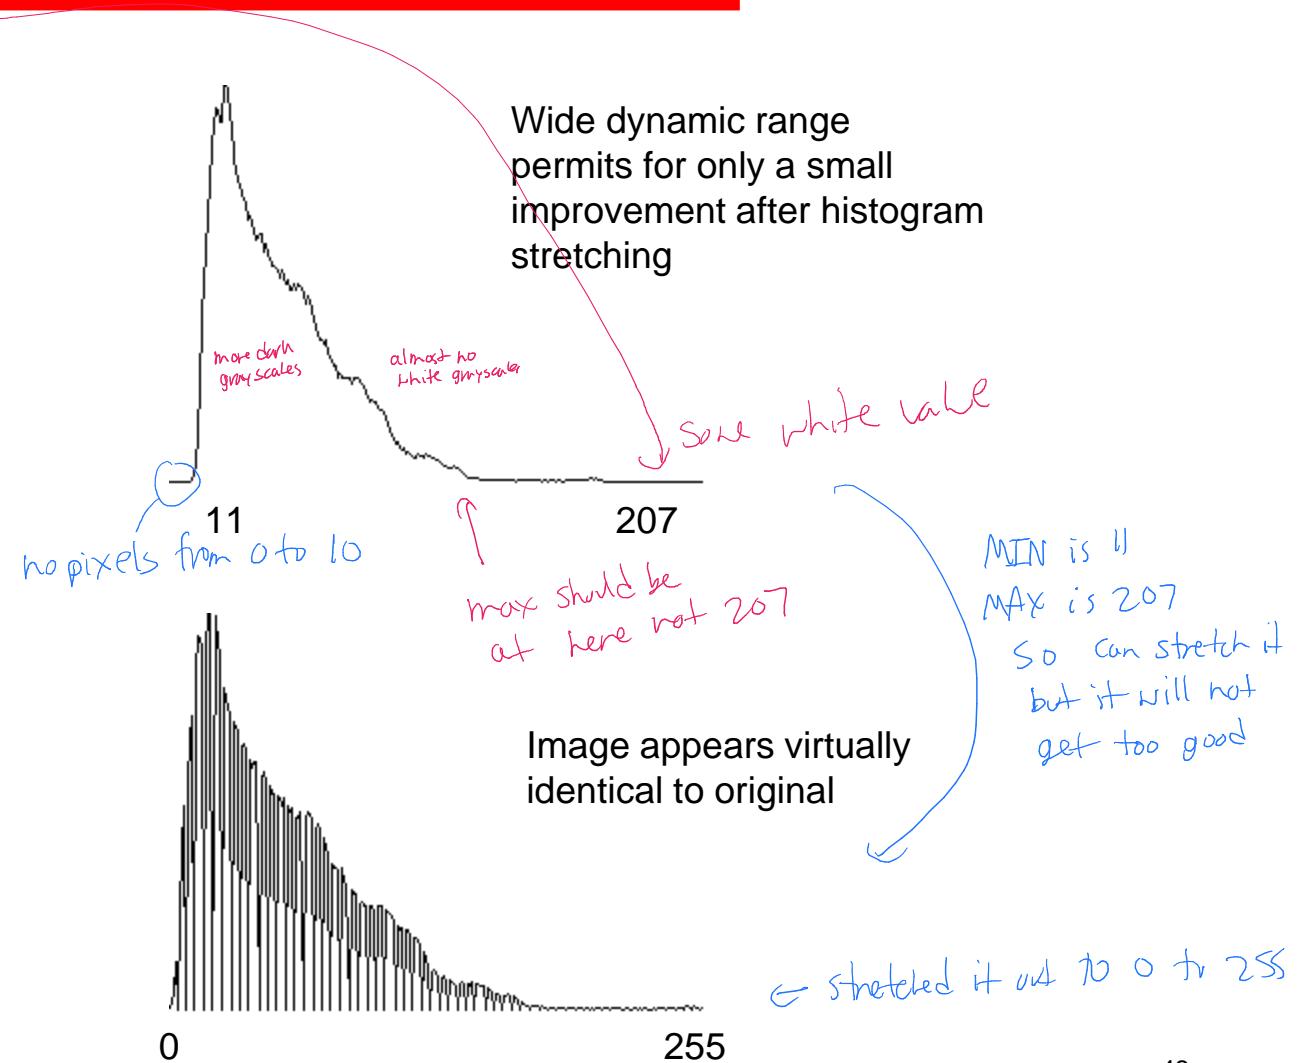
3) Rescale to $[0, 255]$ range

1) Slide histogram down to 0

$$g = \frac{255(f - \text{MIN})}{\text{MAX} - \text{MIN}}$$

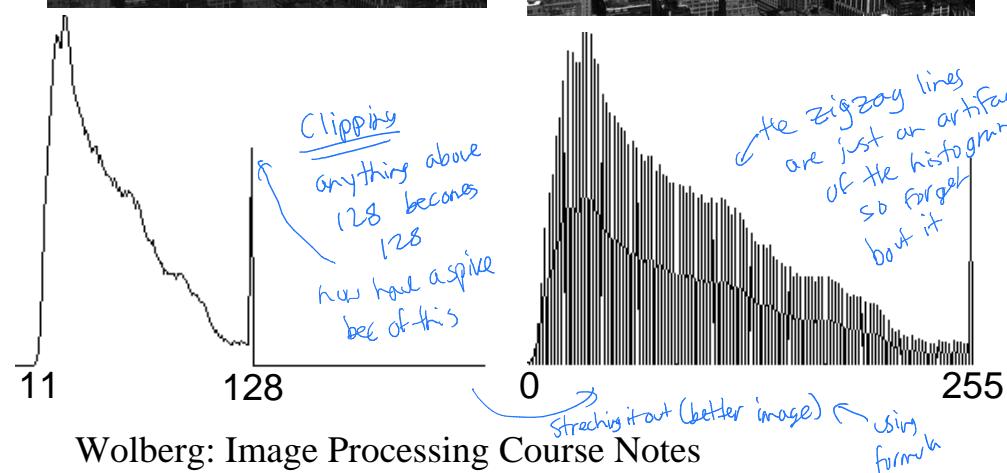
2) Normalize histogram to $[0, 1]$ range

Example (1)



Example (2)

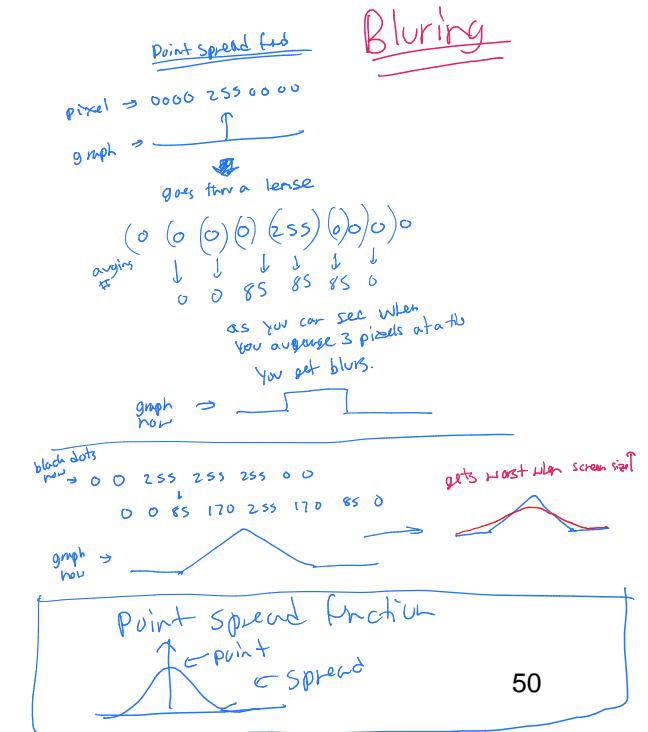
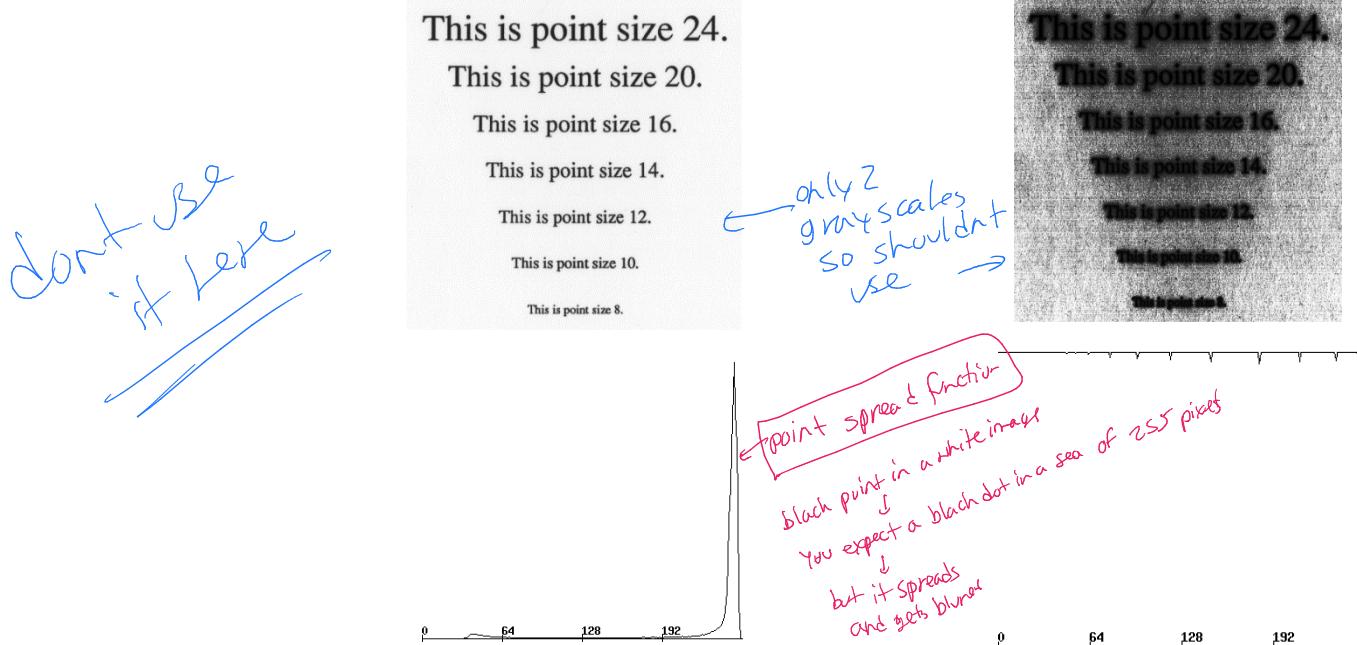
- Improve effectiveness of histogram stretching by clipping intensities first



Flat histogram: every graylevel
is equally present in image

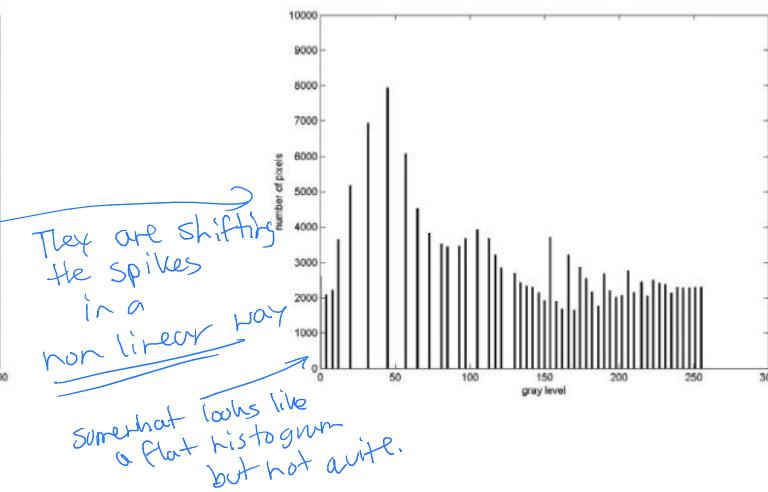
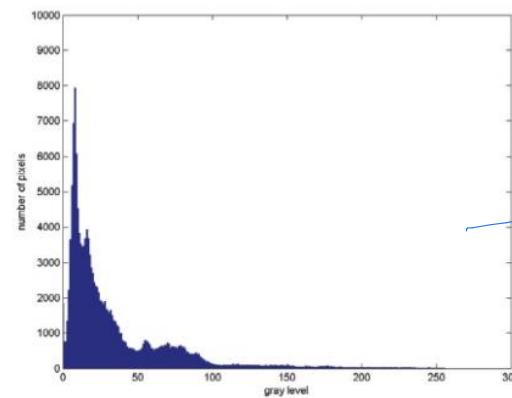
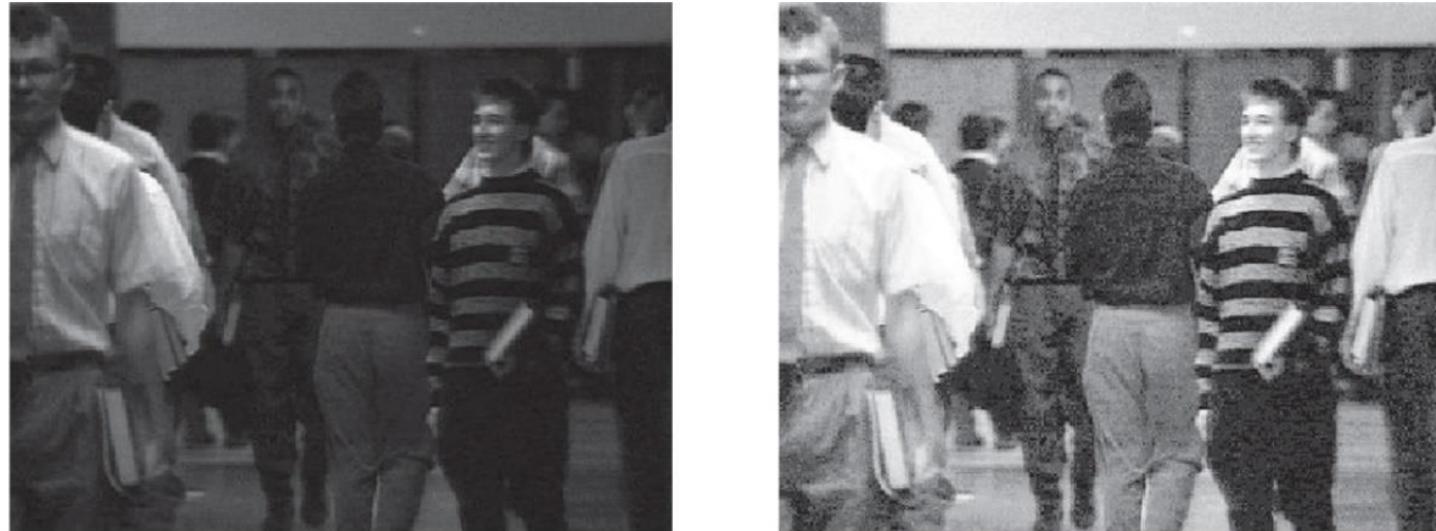
Histogram Equalization

- Produce image with flat histogram
- All graylevels are equally likely
- Appropriate for images with wide range of graylevels
- Inappropriate for images with few graylevels (see below)



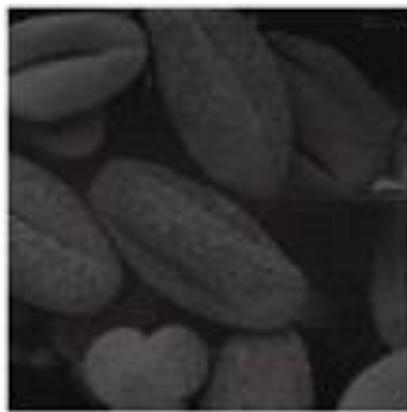
Example

Figure 3.21 The result of histogram equalization applied to an image. The increase in contrast is noticeable. The normalized histogram of the result is much flatter than the original histogram, but it is not completely flat due to discretization effects. Source: Movie *Hoop Dreams*.



Example (1)

before



after



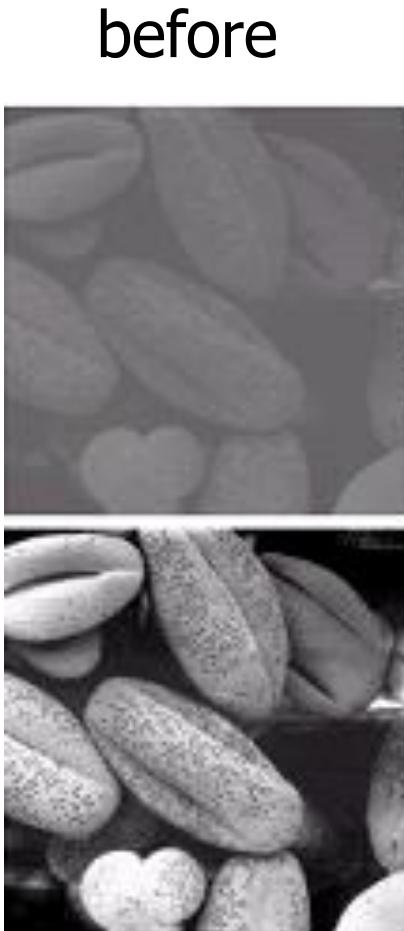
Histogram
equalization



wide dynamic
ranges in
the histogram



Example (2)



Histogram
equalization

The quality is
not improved
much because
the original
image already
has a wide
graylevel scale

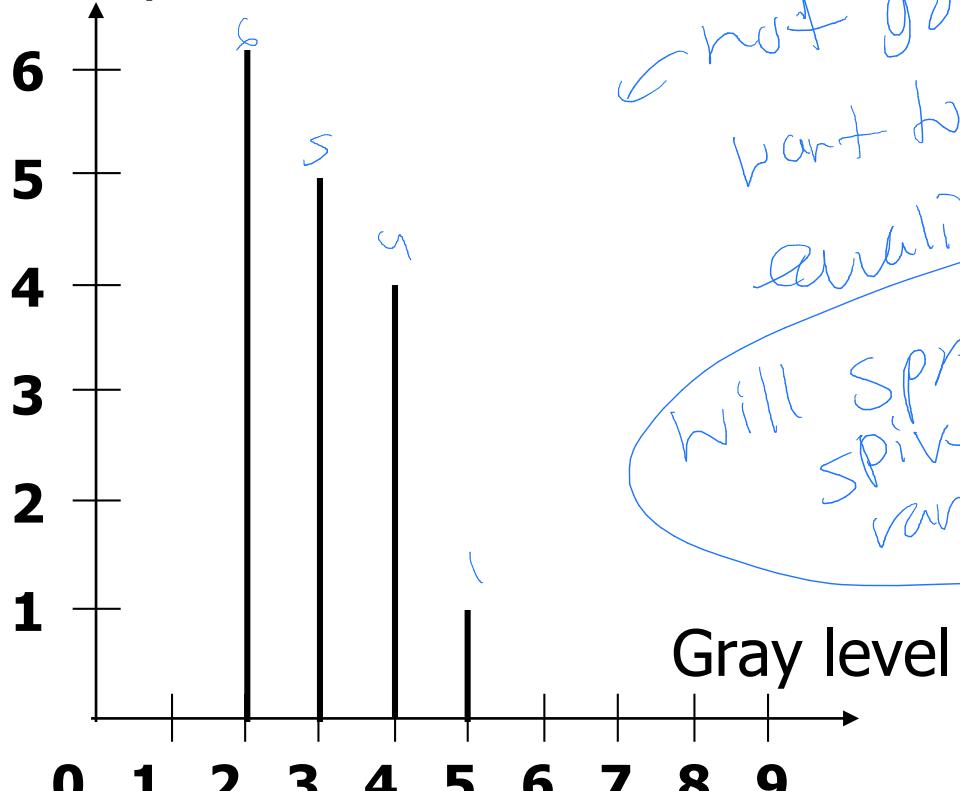
Implementation (1)

2	3	3	2
4	2	4	3
3	2	3	5
2	4	2	4

4x4 image

Gray scale = [0,9]

No. of pixels



histogram

not good
part b
analyze it
will spread spives over the range of 0 to n

*Spreading out
the spines*

Implementation (2)

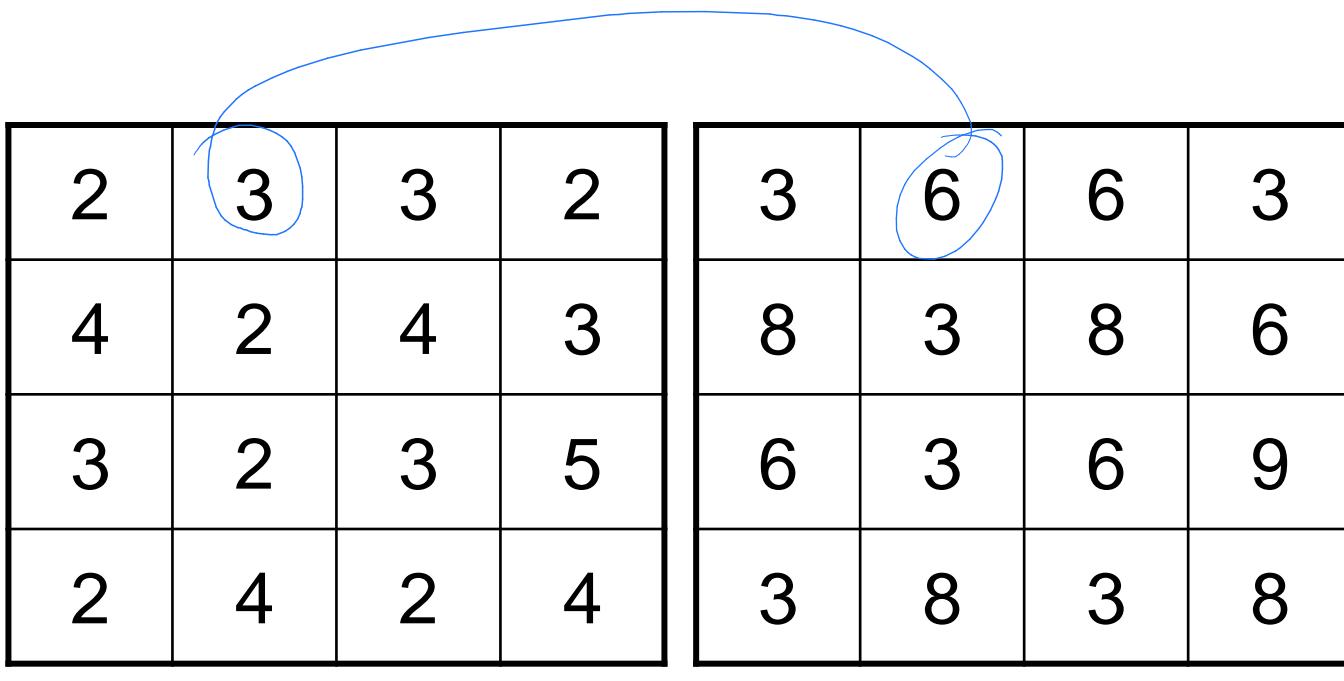
This is a look up table. ("gray val maps to 8^n")

Gray Level(j)	0	1	2	3	4	5	6	7	8	9
No. of pixels	0	0	6	5	4	1	0	0	0	0
$\sum_{j=0}^k n_j$	0	0	6	11	15	16	16	16	16	16
Cumulative Distribution Function (CDF)	$s = \sum_{j=0}^k \frac{n_j}{n}$	0	0	6/16	11/16	15/16	16/16	16/16	16/16	16/16
$s \times 9$	0	0	3.3	6.1	8.4	9	9	9	9	9

every pixel in the original that had a gray val of 2 becomes 3

mapping to the levels

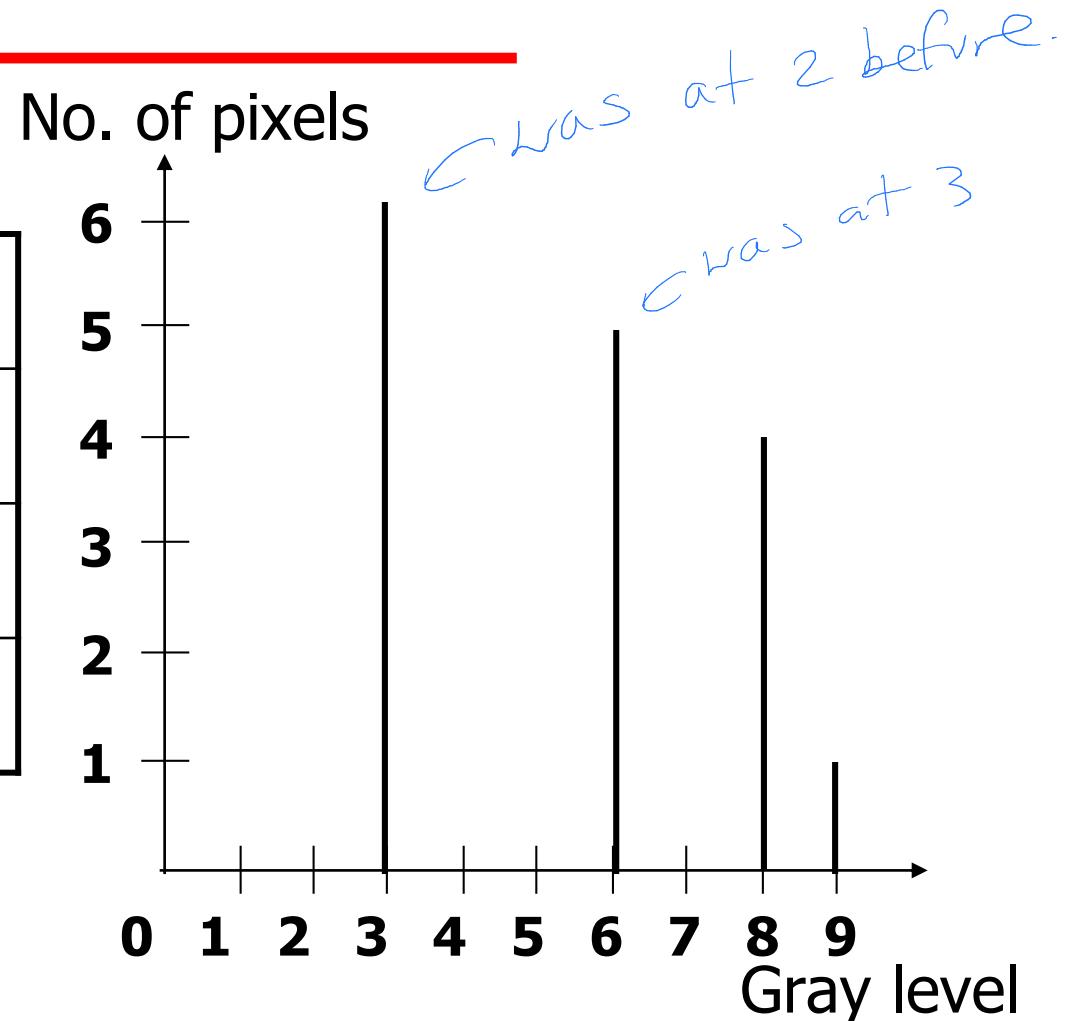
Implementation (3)



Input image

Output image

Gray scale = [0,9]



Histogram equalization

Pseudocode

if we add up all the hist entries
it adds up to 1

ALGORITHM 3.9 Perform histogram equalization on an image

HISTOGRAMEQUALIZE(I)

Input: grayscale image I

Output: histogram-equalized grayscale image I' with increased contrast

- 1 $\bar{h} \leftarrow \text{COMPUTENORMALIZEDHISTOGRAM}(I)$
- 2 $\bar{c} \leftarrow \text{RUNNINGSUM}(\bar{h})$
- 3 **for** $(x, y) \in I$ **do**
- 4 $I'(x, y) \leftarrow \text{ROUND}(255 * \bar{c}[I(x, y)])$
- 5 **return** I'

RUNNINGSUM(a)

Input: 1D array a of $length$ values

Output: 1D running sum s of array

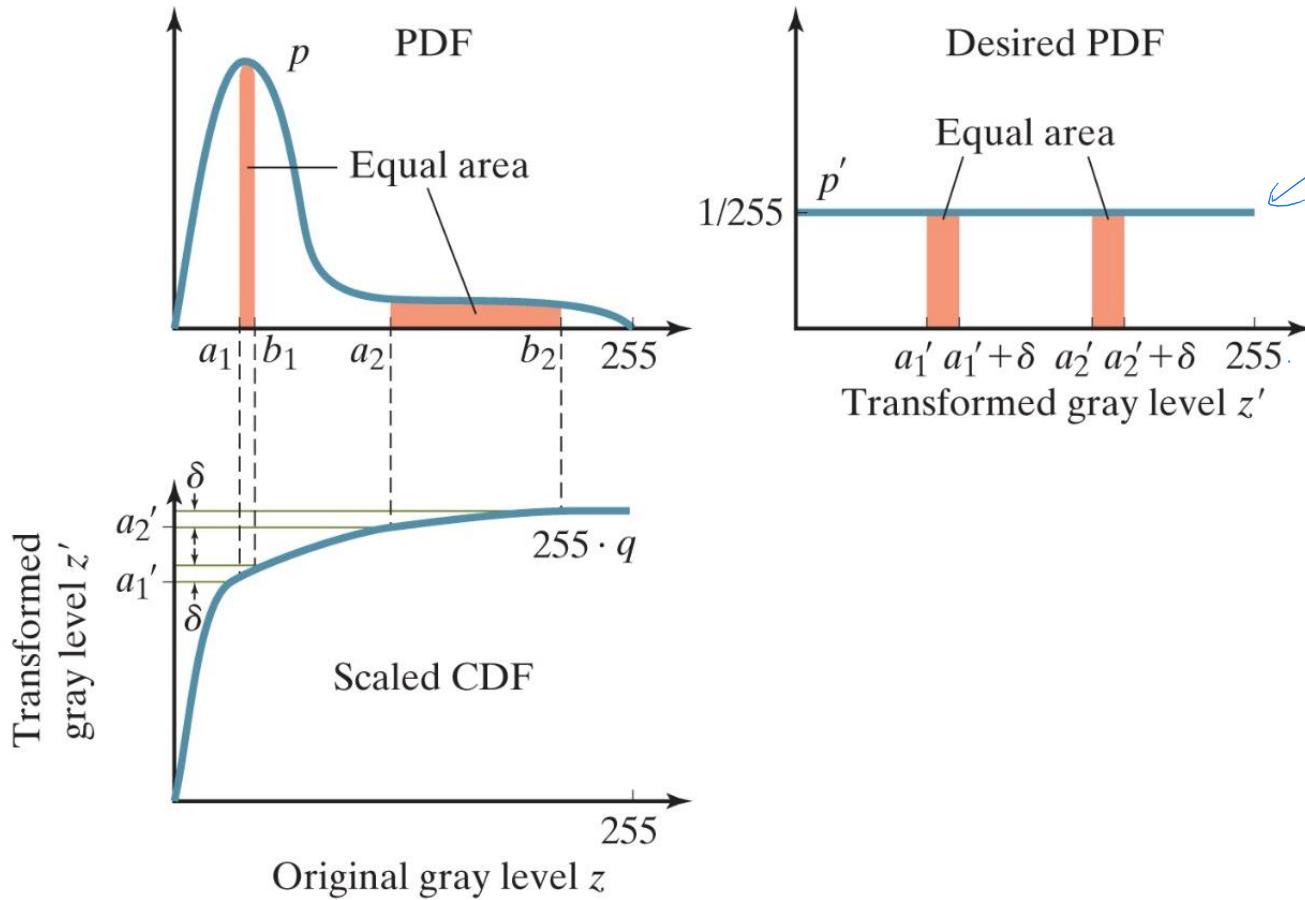
- 1 $s[0] \leftarrow a[0]$
- 2 **for** $k \leftarrow 1$ **to** $length - 1$ **do**
- 3 $s[k] \leftarrow s[k - 1] + a[k]$
- 4 **return** s

basically
going:

$$\sum_{j=0}^k r_j$$

Why It Works

Figure 3.22 Why histogram equalization works. In this example, the histogram of the original image is heavily weighted toward darker pixels. If we let the CDF be the mapping from the old gray level to the new one, the new PDF is flat and therefore weights all gray levels equally. This is because any interval of width δ in the new histogram captures the same number ($\delta/255$) of pixels in the original image. In this example the area within each orange region is identical. Note that discretization effects have been ignored for this illustration.



desired
prob density func:
where every
gray val is
equally likely

Note (1)

- Histogram equalization distributes the graylevels to reach maximum gray (white) because the cumulative distribution function equals 1 when $0 \leq r \leq L-1$
- If $\sum_{j=0}^k n_j$ is slightly different among consecutive k , those graylevels will be mapped to (nearly) identical values as we have to produce an integer grayvalue as output
- Thus, the discrete transformation function cannot guarantee a one-to-one mapping

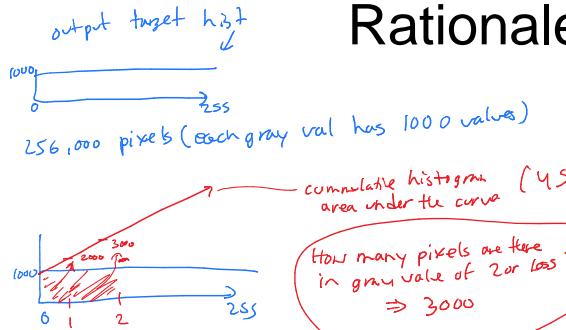
Note (2)

- The implementation given above is widely interpreted as histogram equalization.
- It is readily implemented with a LUT.
- It does not produce a strictly flat histogram.
- There is a more accurate solution. However, it may require a one-to-many mapping that cannot be implemented with a LUT.

Histogram Equalization Objective

Objective: we want a uniform histogram.

Rationale: maximize image entropy.



$$h_1(v_{out}) = \text{constant} = \frac{\text{total}}{MXGRAY}$$

value of the pixel in output image

the value of the hist should be this

$$= h_{avg}$$

$$c_1(v_{out}) = (v_{out} + 1) * h_{avg}$$

This is a special case of histogram matching.

Perfectly flat histogram: $H[i] = \text{total}/MXGRAY$ for $0 \leq i < MXGRAY$.

If $H[v] = k * h_{avg}$ then v must be mapped onto k different levels, from v_1 to v_k . This is a one-to many mapping.

a hist where ist just constant.

256 gray vals
1000 pixels in image
↓ sum it up: 256,000 vals
each gray value has 1000 pixels???

Histogram Equalization Mappings

2.26
this and last slide
be explained

If g_my lvl o want to map to 0,1,2. \rightarrow map it to 1

Rule 1: Always map v onto $(v_1 + v_k)/2$. (This does not result in a flat histogram, but one where brightness levels are spaced apart).

Rule 2: Assign at random one of the levels in $[v_1, v_k]$. This can result in a loss of contrast if the original histogram had two distinct peaks that were far apart (i.e., an image of text).

Rule 3: Examine neighborhood of pixel, and assign it a level from $[v_1, v_k]$ which is closest to neighborhood average. This can result in blurriness; more complex.

Rule (1) creates a lookup table beforehand.

ex 2

Rules (2) and (3) are runtime operations.

Better Implementation (1)

```
void histeq(imageP I1, imageP I2) ← pointer to image
{
    int i, R;
    int left[MXGRAY], width[MXGRAY];
    uchar *in, *out;
    long total, Hsum, Havg, histo[MXGRAY];
    /* assume dimensions are the same */

    /* total number of pixels in image */
    total = (long) I1->width * I1->height;

    /* init I2 dimensions and buffer */
    I2->width = I1->width;
    I2->height = I1->height;
    I2->image = (uchar *) malloc(total);

    /* init input and output pointers */
    in = I1->image;                                /* input image buffer */
    out = I2->image;                               /* output image buffer */

    /* compute histogram */
    for(i=0; i<MXGRAY; i++) histo[i] = 0; /* clear histogram */
    for(i=0; i<total; i++) histo[in[i]]++; /* eval histogram */
    R = 0;
    Hsum = 0;
    Havg = total / MXGRAY;
    /* right end of interval */
    /* cumulative value for interval */
    /* interval value for uniform histogram */
}
```

pointer to image
can access height, width
histo gun
equalize

but guaranteed

set all stuff in hist to 0
values from image

running all the pixels in input image.

What is histo[...] + the??!

Accumulate avg for integral

Better Implementation (2)

ends at
2:37

```
/* evaluate remapping of all input gray levels;
 * Each input gray value maps to an interval of valid output values.
 * The endpoints of the intervals are left[] and left[]+width[].
 */
for(i=0; i<MXGRAY; i++) {
    left[i] = R;          /* maps to 0 */
    Hsum += histo[i];    /* accumulate sum history → running sum + current histo value */
    while(Hsum>Havg && R<MXGRAY-1) { /* make interval wider */
        Hsum -= Havg;      /* adjust Hsum */
        R++;               /* update right end */
    }
    width[i] = R - left[i] + 1; /* width of interval */
}

/* visit all input pixels and remap intensities */
for(i=0; i<total; i++) {
    if(width[in[i]] == 1) out[i] = left[in[i]];
    else {
        /* in[i] spills over into width[] possible values */
        /* randomly pick from 0 to width[i] */
        R = ((rand()&0x7fff)*width[in[i]])>>15; /* 0 <= R < width */
        out[i] = left[in[i]] + R;
    }
}
```

for every gray val mapped to some width (interval)

maps and out gray vals?

remap to interval

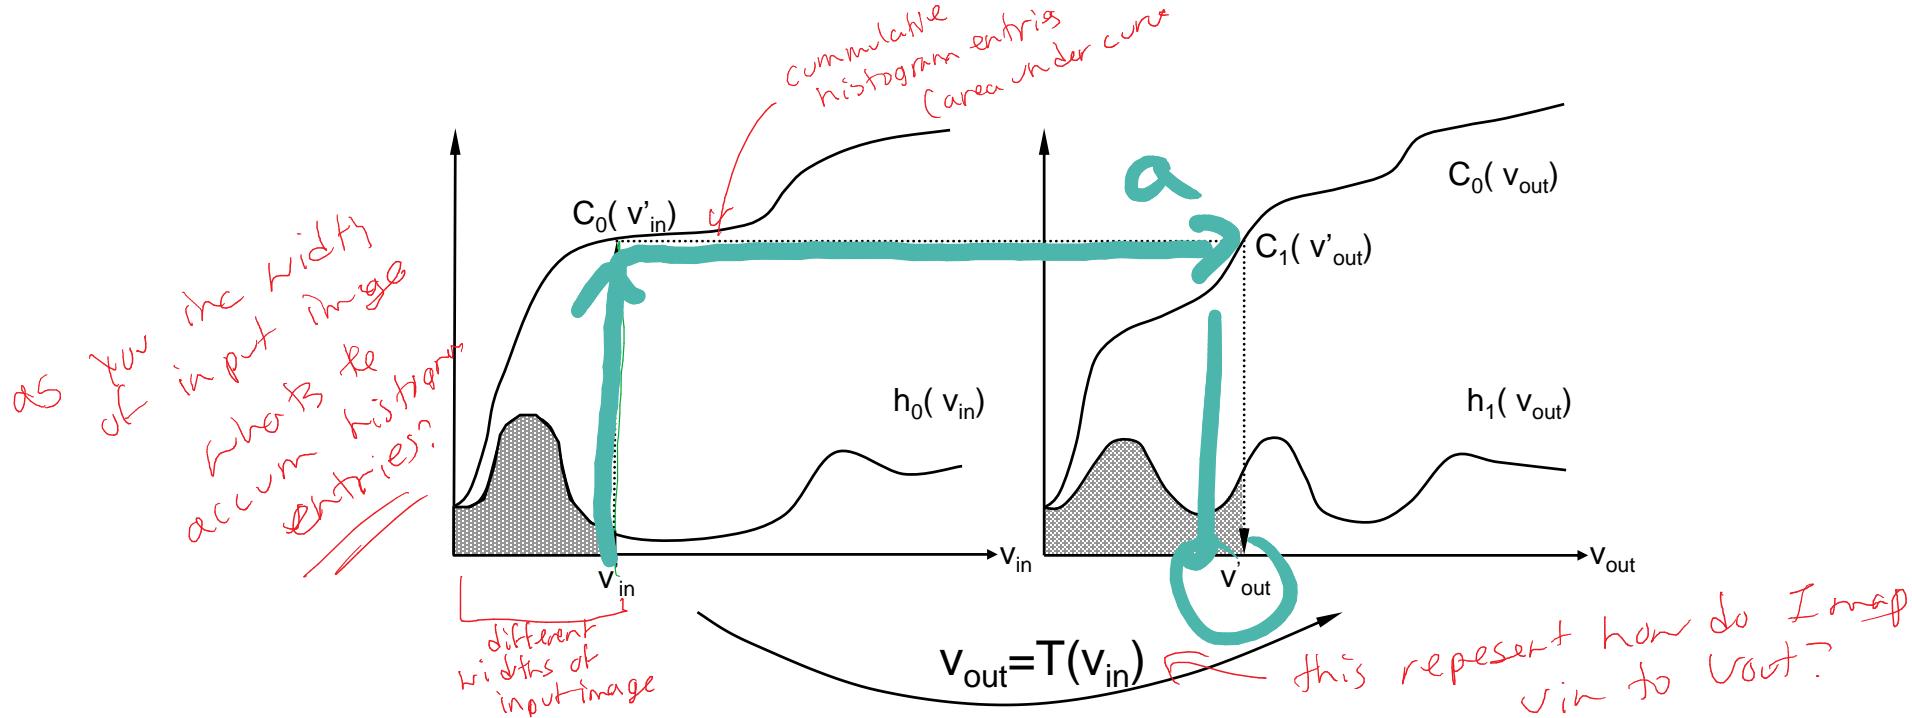
300U 1000

pick a random to add it to the left interval

Note

- Histogram equalization has a disadvantage:
it can generate only one type of output image.
- With histogram specification we can specify the shape of the histogram that we wish the output image to have.
- It doesn't have to be a uniform histogram.
- Histogram specification is a trial-and-error process.
- There are no rules for specifying histograms, and one must resort to analysis on a case-by-case basis for any given enhancement task.

Histogram Matching



In the figure above, $h()$ refers to the histogram, and $c()$ refers to its cumulative histogram. Function $c()$ is a *monotonically increasing function defined as:*

$$c(v) = \int_0^v h(u) du$$

always rises
but never falls
bec (-area) don't exist.

Histogram Matching Rule

Let $v_{out} = T(v_{in})$ If $T()$ is a unique, monotonic function then

$$\int_0^{v_{out}} h_1(u) du = \int_0^{v_{in}} h_0(u) du$$

This can be restated in terms of the histogram matching rule:

$$c_1(v_{out}) = c_0(v_{in})$$

Where $c_1(v_{out}) = \# \text{ pixels } \leq v_{out}$, and $c_0(v_{in}) = \# \text{ pixels } \leq v_{in}$.

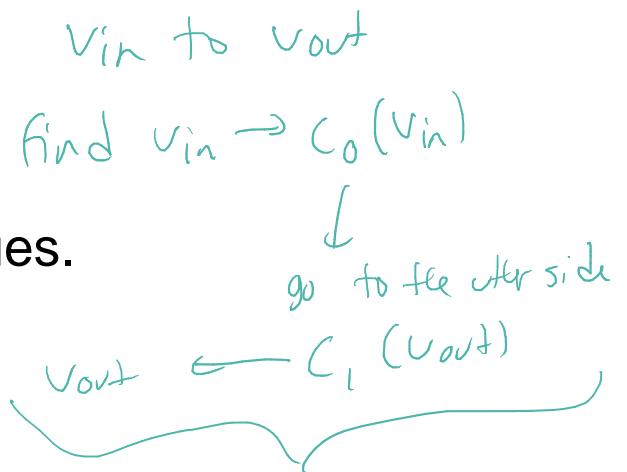
This requires that

$$v_{out} = c_1^{-1}(c_0(v_{in}))$$

which is the basic equation for histogram matching techniques.

$$c_1(v_{out}) = c_0(v_{in})$$

$$v_{out} = c_1^{-1}(c_0(v_{in}))$$



2 : WS
starts

Pseudocode

ALGORITHM 3.10 Perform histogram matching on an image

HISTOGRAMMATCH(I, h_{ref})

Input: grayscale image I , reference normalized graylevel histogram \bar{h}_{ref}
Output: grayscale image I' whose normalized graylevel histogram closely matches \bar{h}_{ref}

```
1   $\bar{h} \leftarrow \text{COMPUTENORMALIZEDHISTOGRAM}(I)$                                 ➤ Compute the normalized histogram of the image,  
2   $\bar{c} \leftarrow \text{RUNNINGSUM}(\bar{h})$                                          then compute the CDF of the image,  
3   $\bar{c}_{ref} \leftarrow \text{RUNNINGSUM}(\bar{h}_{ref})$                                      as well as the desired CDF.  
4  for  $\ell \leftarrow 0$  to 255 do  
5     $\ell' \leftarrow 255$   
6    repeat  
7       $f[\ell] \leftarrow \ell'$   
8       $\ell' \leftarrow -1$   
9      while  $\ell' \geq 0$  AND  $\bar{c}_{ref}[\ell'] > \bar{c}[\ell]$   
10     for  $(x, y) \in I$  do  
11        $I'(x, y) \leftarrow f[I(x, y)]$   
12     return  $I'$   
13  
14
```

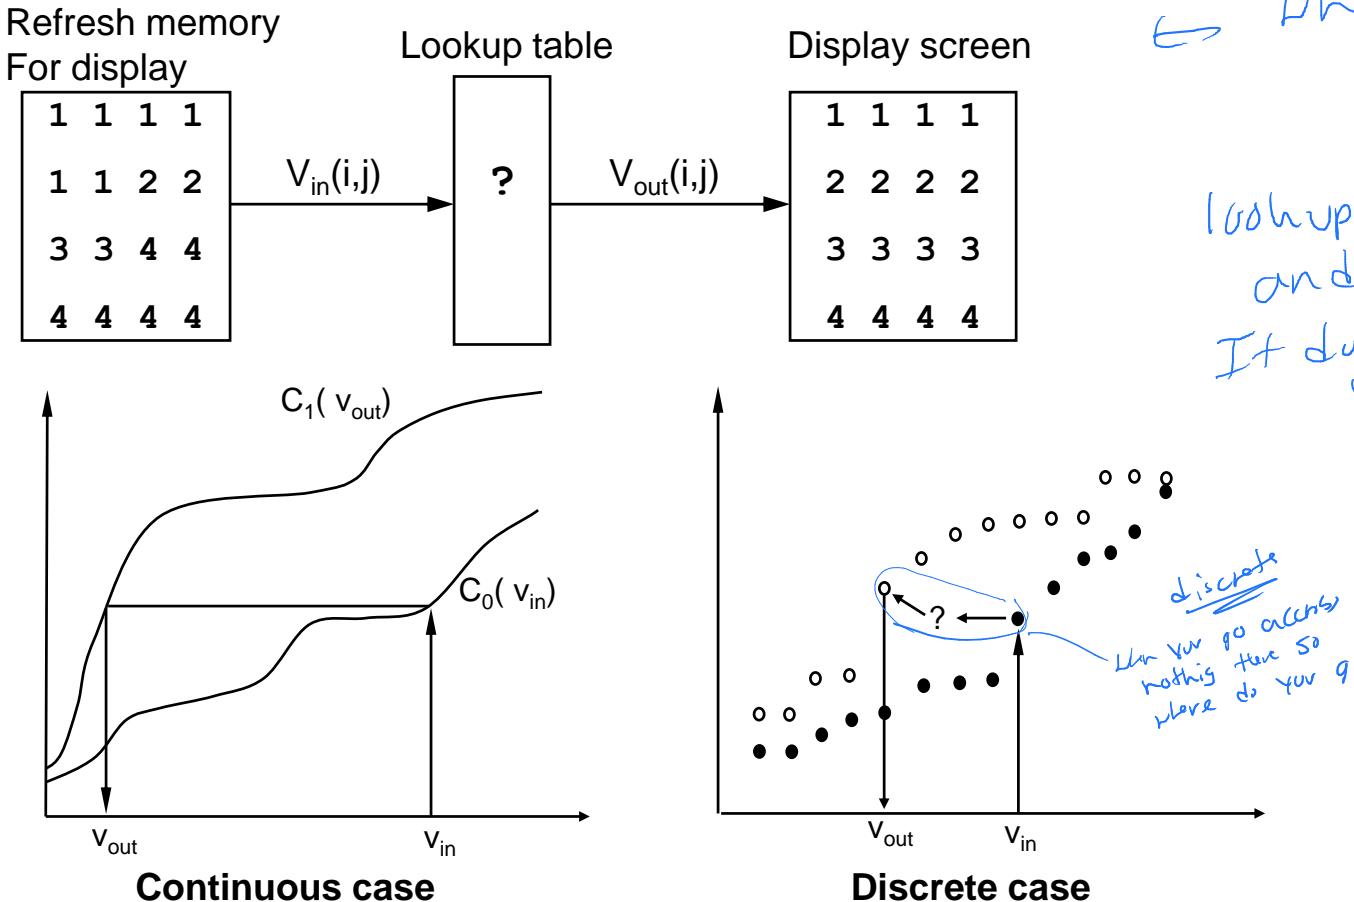
➤ For each possible gray level ℓ , set $f[\ell] \approx \bar{c}_{ref}^{-1}[\bar{c}(\ell)]$.
This is done by finding ℓ' such that $\bar{c}_{ref}[\ell'] \approx \bar{c}[\ell]$,
and setting $f[\ell] = \ell'$. To handle discretization effects,
 ℓ' is set to the maximum possible gray level,
then repeatedly decremented
until $\bar{c}_{ref}[\ell'] \leq \bar{c}[\ell]$.

➤ Once the mapping f has been determined,
it is applied to all pixels in the image.

2.4)

Histograms are Discrete

- Impossible to match all histogram pairs because they are discrete.



Problems with Discrete Case

- The set of input pixel values is a discrete set, and all the pixels of a given value are mapped to the same output value. For example, all six pixels of value one are mapped to the same value so it is impossible to have only four corresponding output pixels.
- No inverse for c_1 in $v_{out} = c_1^{-1}(c_0(v_{in}))$ because of discrete domain. Solution: choose v_{out} for which $c_1(v_{out})$ is closest to $c_0(v_{in})$.
- $v_{in} \rightarrow v_{out}$ such that $|c_1(v_{out}) - c_0(v_{in})|$ is a minimum

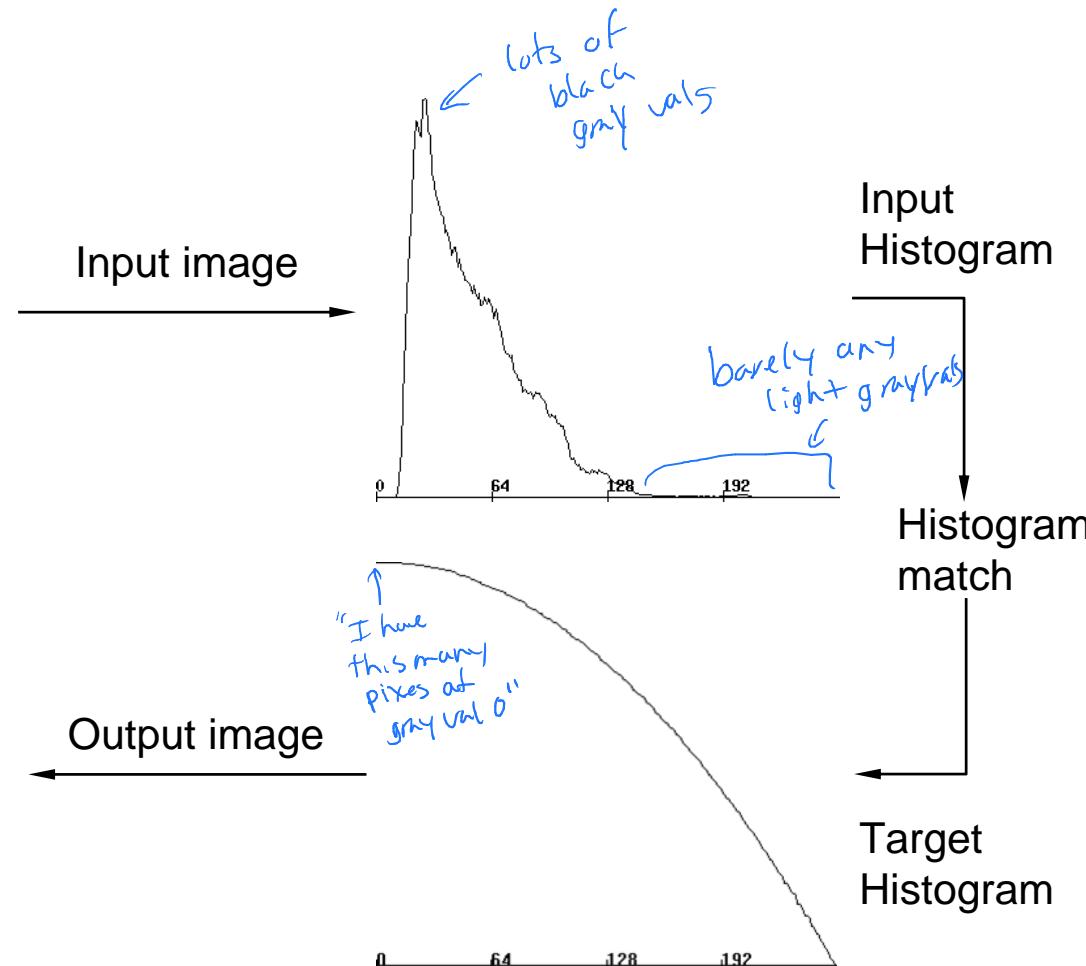
Histogram Matching Example (1)



Input image



Output image



Pixel#
 2 500 500 500 pixels
gray vals
 0 → 0, 1, 2, 3
 1 → 3, 4, 5, 6
 2 → 6

how much pixels the gray val picks up (limit 500)
 but say wants to make
 40 pixels to 3

500-40 need to know that only got 500 - 40 remaining

need to know that gray val 2 needs to share to
 $\frac{500-20}{500-70}$

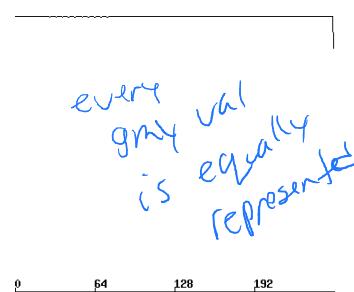
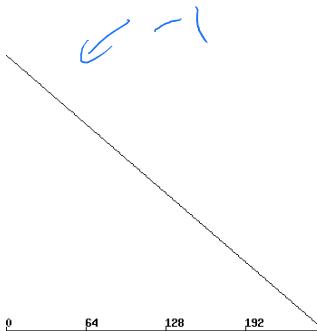
Make a new array
 that tells you
 how much is leftover,
 from previous pixel

3:0^q ends

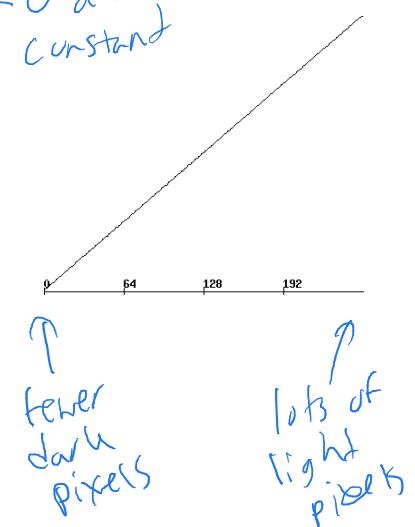
Histogram Matching Example (2)

If I apply the code on the next slide, these will be spiky.

3:00 to 2:03 we explain what extra code we need to figure out for the HTH to work



↙ 0 all constant



Implementation (1)

```

int histogramMatch(imageP I1, imageP histo, imageP I2)
{
    int i, p, R;
    int left[MXGRAY], right[MXGRAY];
    int total, Hsum, Havg, h1[MXGRAY], *h2;
    unsigned char *in, *out;
    double scale;

    /* total number of pixels in image */
    total = (long) I1->height * I1->width;

    /* init I2 dimensions and buffer */
    I2->width = I1->width;
    I2->height = I1->height;
    I2->image = (unsigned char *) malloc(total);

    in = I1->image;                      /* input image buffer */
    out = I2->image;                     /* output image buffer */

    for(i=0; i<MXGRAY; i++) h1[i] = 0; /* clear histogram */
    for(i=0; i<total; i++) h1[in[i]]++; /* eval histogram */

```

input image *output ↗*

*output histogram
(an array of vals)
ex flat hist: [3, 3, 3, ...]*

Implementation (2)

```
/* target histogram */  
h2 = (int *) histo->image;  
  
/* normalize h2 to conform with dimensions of I1 */  
for(i=Havg=0; i<MXGRAY; i++) Havg += h2[i];  
scale = (double) total / Havg;          total # pixel in image  
if(scale != 1) for(i=0; i<MXGRAY; i++) h2[i] *= scale;  
  
R = 0;  
Hsum = 0;  
/* evaluate remapping of all input gray levels;  
   Each input gray value maps to an interval of valid output values.  
   The endpoints of the intervals are left[] and right[] */  
for(i=0; i<MXGRAY; i++) {  
    left[i] = R;                      /* left end of interval */  
    Hsum += h1[i];                    /* cumulative value for interval */  
    while(Hsum>h2[R] && R<MXGRAY-1) { /* compute width of interval */  
        Hsum -= h2[R];                /* adjust Hsum as interval widens */  
        R++;                         /* update */  
    }  
    right[i] = R;                    /* init right end of interval */  
}
```

histogram is just a curve. don't know size of image. This fixes that

*loop thru i^2
0 to 255*

*add up all vals of h2 histogram
why? may not know the total # of pixels of image
total # of hist entries = to total # of pixels in image*

ends
3.0\

Implementation (3)

```
/* clear h1 and reuse it below */  
for(i=0; i<MXGRAY; i++) h1[i] = 0;  
  
/* visit all input pixels */  
for(i=0; i<total; i++) {  
    p = left[in[i]];  
    if(h1[p] < h2[p]) /* mapping satisfies h2 */  
        out[i] = p;  
    else  
        out[i] = p = left[in[i]] = MIN(p+1, right[in[i]]);  
    h1[p]++;  
}
```

new left
of the
interval

h_1 = running counter of how
many

h_2 = output histogram normalized

Local Pixel Value Mappings

- Histogram processing methods are global, in the sense that pixels are modified by a transformation function based on the graylevel content of an entire image.
- We sometimes need to enhance details over small areas in an image, which is called a local enhancement.
- Solution: apply transformation functions based on graylevel distribution within pixel neighborhood.

General Procedure

- Define a square or rectangular neighborhood.
- Move the center of this area from pixel to pixel.
- At each location, the histogram of the points in the neighborhood is computed and histogram equalization, histogram matching, or other graylevel mapping is performed.
- Exploit easy histogram update since only one new row or column of neighborhood changes during pixel-to-pixel translation.
- Another approach used to reduce computation is to utilize nonoverlapping regions, but this usually produces an undesirable checkerboard effect.

Example: Local Enhancement

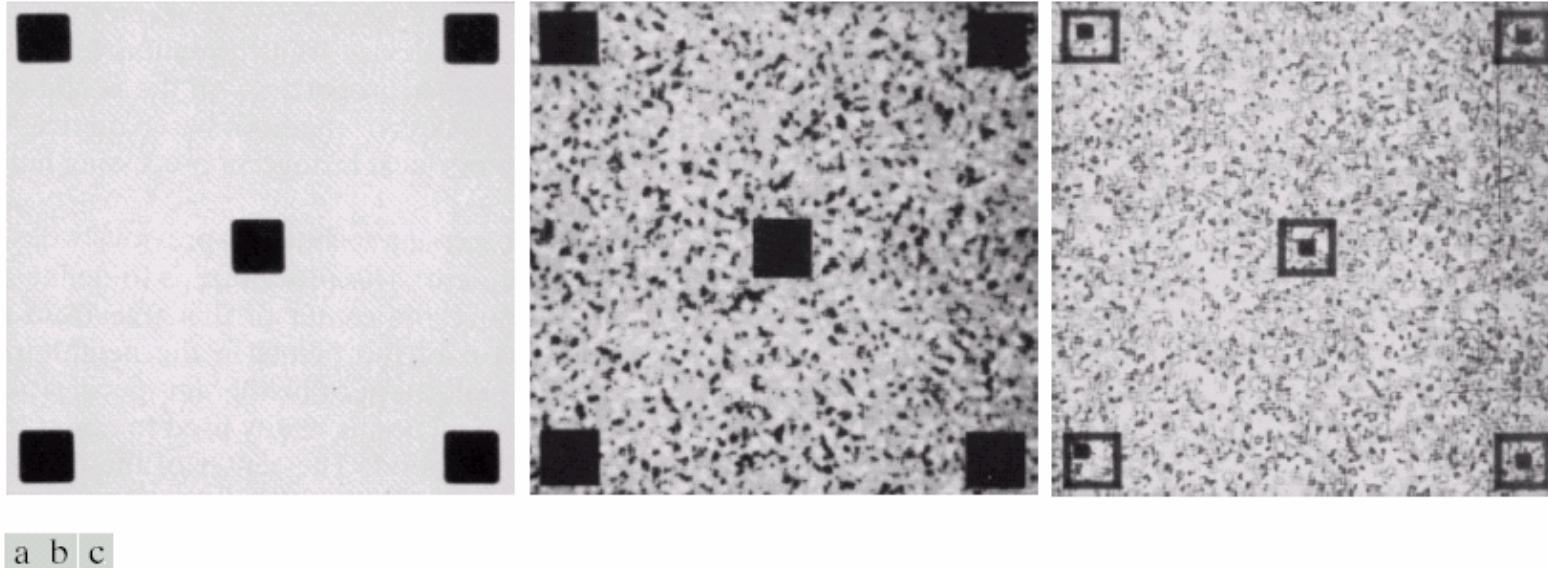


FIGURE 3.23 (a) Original image. (b) Result of global histogram equalization. (c) Result of local histogram equalization using a 7×7 neighborhood about each pixel.

- a) Original image (slightly blurred to reduce noise)
- b) global histogram equalization enhances noise & slightly increases contrast but the structural details are unchanged
- c) local histogram equalization using 7×7 neighborhood reveals the small squares inside of the larger ones in the original image.

3:09
start

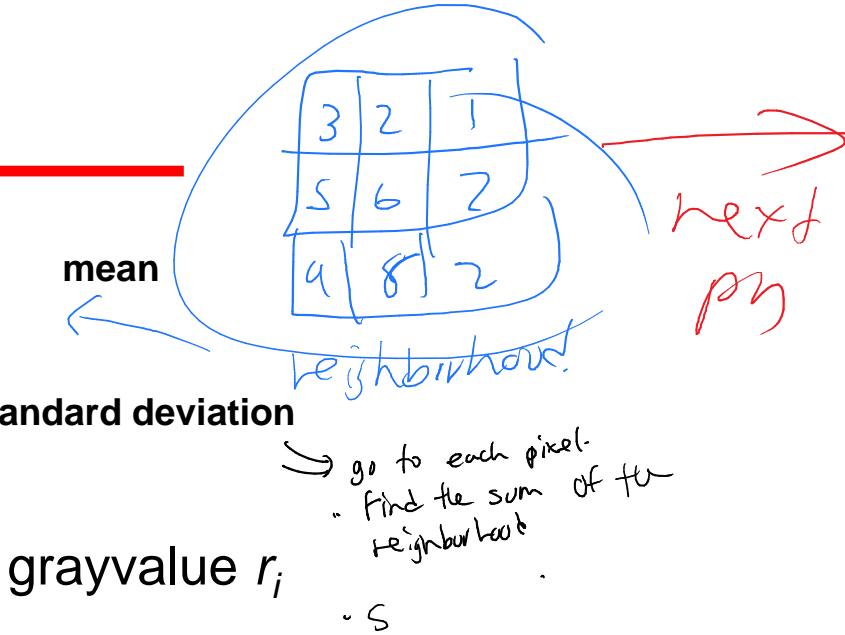
Definitions (1)

$$\mu(x, y) = \frac{1}{n} \sum_{i,j} f(i, j) \quad \text{Find average value of the pixel neighborhood}$$

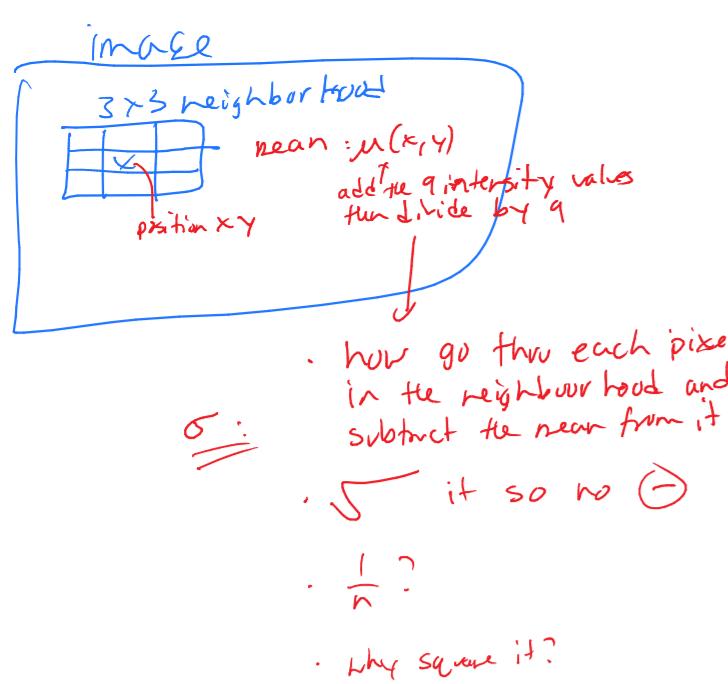
$$\sigma(x, y) = \sqrt{\frac{1}{n} \sum_{i,j} (f(i, j) - \mu(x, y))^2} \quad \text{take value of center of neighborhood}$$

- Let $p(r_i)$ denote the normalized histogram entry for grayvalue r_i , for $0 \leq i < L$ where L is the number of graylevels.
- It is an estimate of the probability of occurrence of graylevel r_i .
- Mean m can be rewritten as

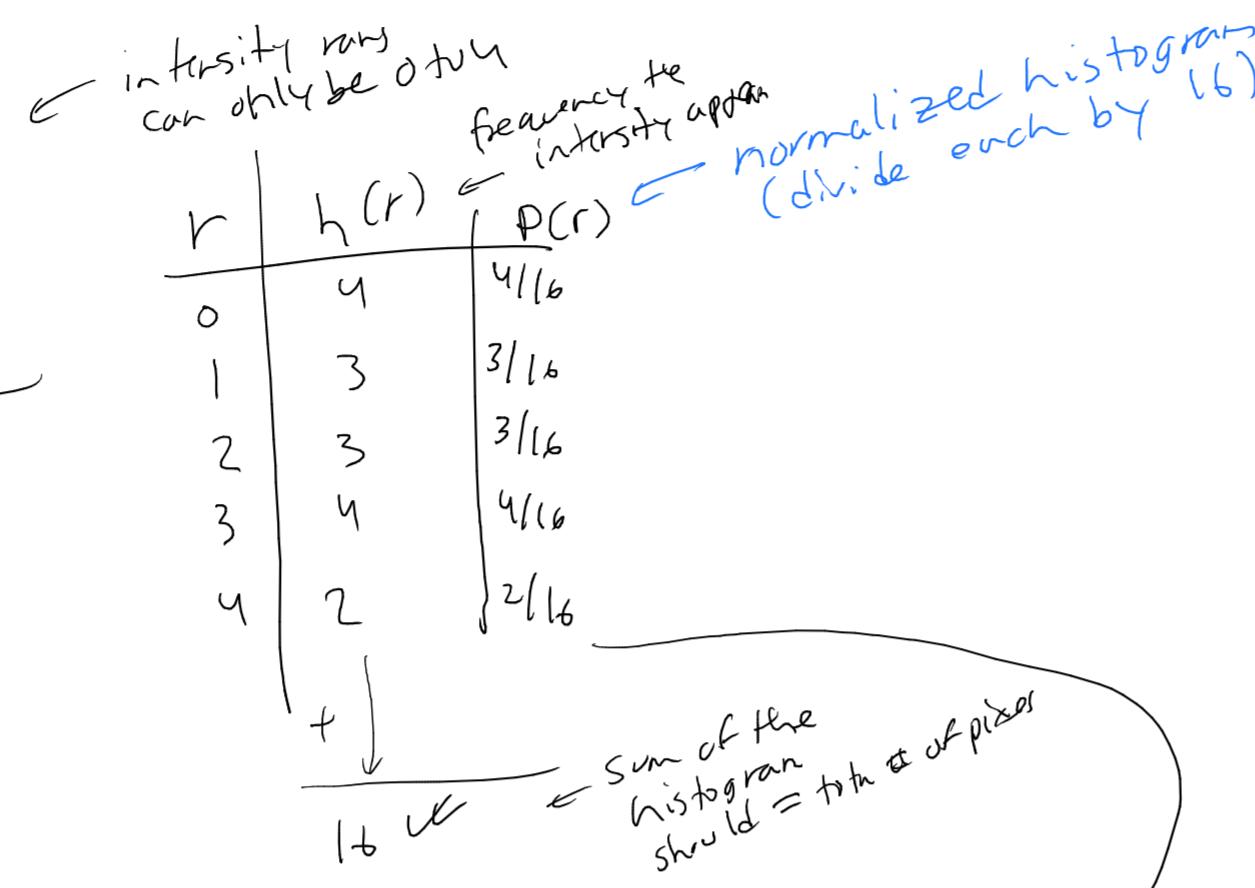
$$m = \sum_{i=0}^{L-1} r_i p(r_i) \quad \begin{array}{l} \text{probabilty of gray level} \\ \text{find the gray val} \\ \text{probabilty of the gray value} \\ \uparrow \\ \text{sum for all poss gray values from 0 to 255} \end{array}$$



0 to 1 be normalized?

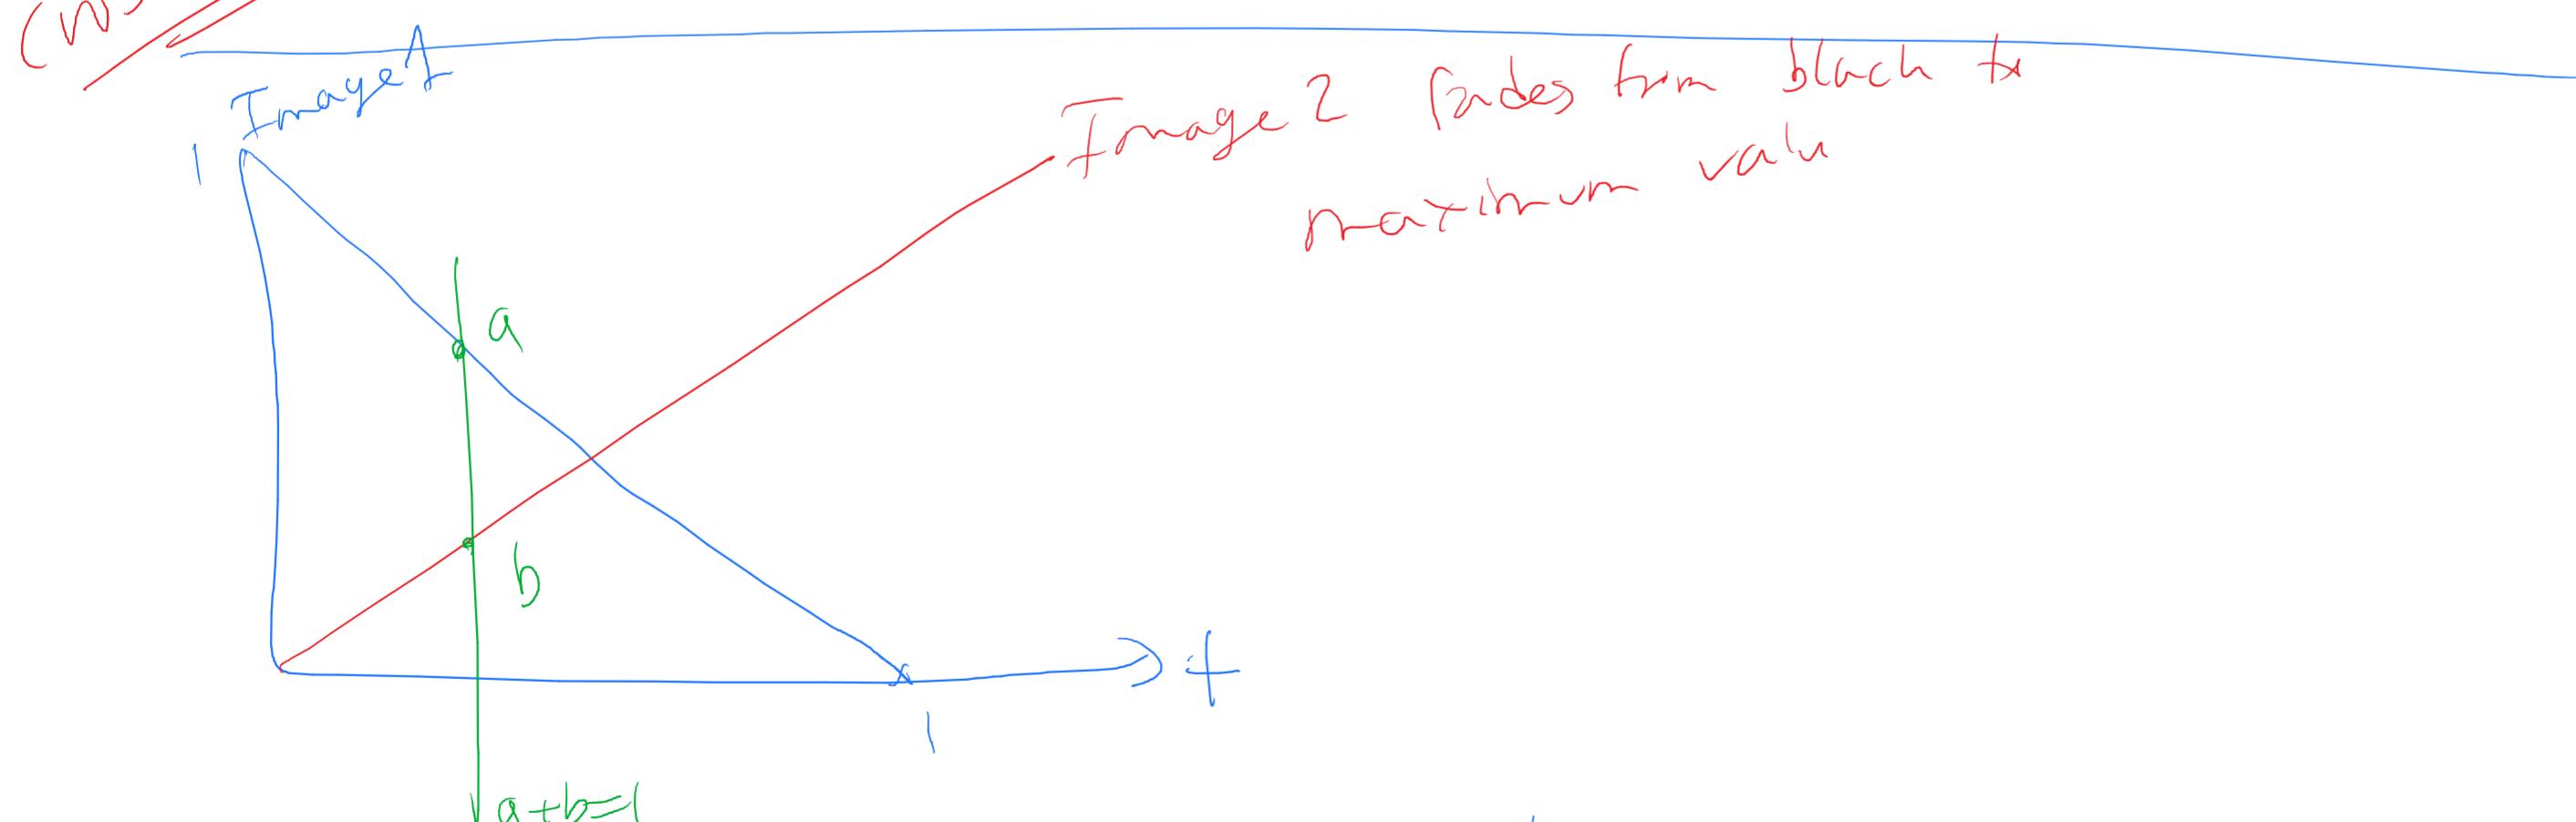


0	0	1	2
3	0	2	1
3	3	3	4
4	0	1	2



$$\begin{aligned}
 \text{mean} &= m = \sum_{i=0}^4 r_i p(r_i) \\
 &= 0 \cdot \frac{4}{16} + 1 \cdot \frac{3}{16} + 2 \cdot \frac{3}{16} + 3 \cdot \frac{4}{16} + 4 \cdot \frac{2}{16} \\
 &\quad \text{or can just add up all the 0's 1's 2's 3's 4's} \\
 &= 0 + 3 + 6 + 12 + 8 \\
 &\quad \text{faster way} = \frac{0 + 3 + 6 + 12 + 8}{16} \quad \text{u's and } \div 16
 \end{aligned}$$

(MSS dissolve) \rightarrow fade out of one
fade in of another



the scale factor applied to both = 1

image 1 fades over time

Definitions (2)

- The nth moment of r about its mean is defined as

$$\mu_n(r) = \sum_{i=0}^{L-1} (r_i - m)^n p(r_i)$$

- It follows that:

$$\mu_0(r) = 1 \quad \text{0th moment}$$

$$\mu_1(r) = 0 \quad \text{1st moment}$$

$$\mu_2(r) = \sum_{i=0}^{L-1} (r_i - m)^2 p(r_i) \quad \text{2nd moment}$$

- The second moment is known as variance $\sigma^2(r)$
- The standard deviation is the square root of the variance.
- The mean and standard deviation are measures of average grayvalue and average contrast, respectively.

Example: Statistical Differencing

- Produces the same contrast throughout the image.
- Stretch $f(x, y)$ away from or towards the local mean to achieve a balanced local standard deviation throughout the image.
- σ_0 is the desired standard deviation and it controls the amount of stretch.
- The local mean can also be adjusted:

$$g(x, y) = \alpha m_0 + (1 - \alpha)\mu(x, y) + (f(x, y) - \mu(x, y)) \frac{\sigma_0}{\sigma(x, y)}$$

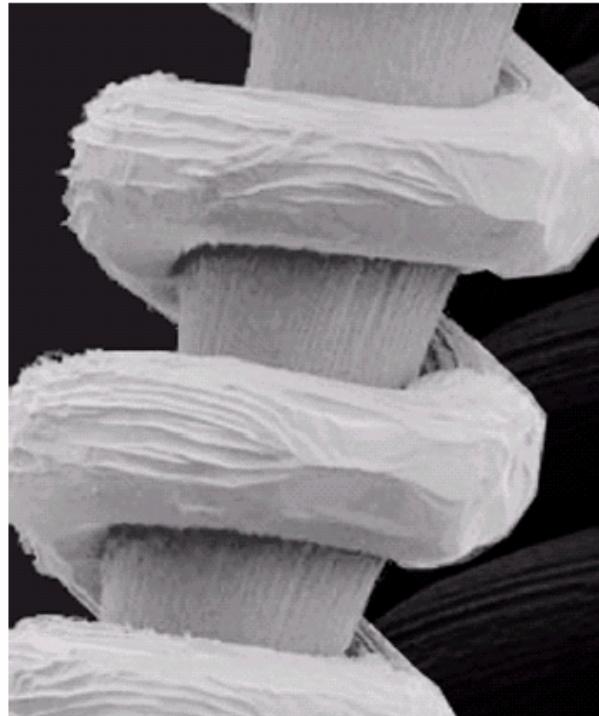
- m_0 is the mean to force locally and α controls the degree to which it is forced.
- To avoid problems when $\sigma(x, y) = 0$,

$$g(x, y) = \alpha m_0 + (1 - \alpha)\mu(x, y) + (f(x, y) - \mu(x, y)) \frac{\beta\sigma_0}{\sigma_0 + \beta\sigma(x, y)}$$

- Speedups can be achieved by dividing the image into blocks (tiles), exactly computing the mean and standard deviation at the center of each block, and then linearly interpolating between blocks in order to compute an approximation at any arbitrary position. In addition, the mean and standard deviation can be computed incrementally.

Example: Local Statistics (1)

FIGURE 3.24 SEM image of a tungsten filament and support, magnified approximately 130×. (Original image courtesy of Mr. Michael Shaffer, Department of Geological Sciences, University of Oregon, Eugene).



The filament in the center is clear.

There is another filament on the right side that is darker and hard to see.

Goal: enhance dark areas while leaving the light areas unchanged.

Example: Local Statistics (2)

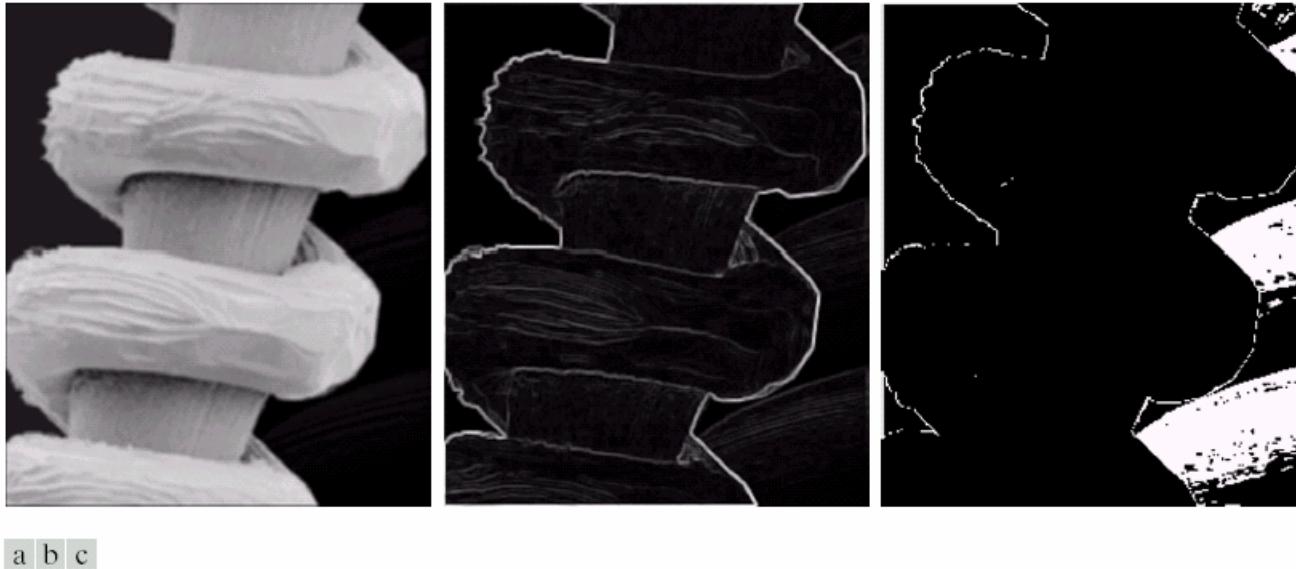


FIGURE 3.25 (a) Image formed from all local means obtained from Fig. 3.24 using Eq. (3.3-21). (b) Image formed from all local standard deviations obtained from Fig. 3.24 using Eq. (3.3-22). (c) Image formed from all multiplication constants used to produce the enhanced image shown in Fig. 3.26.

Solution: Identify candidate pixels to be dark pixels with low contrast.
Dark: local mean $< k_0 * \text{global mean}$, where $0 < k_0 < 1$.
Low contrast: $k_1 * \text{global variance} < \text{local variance} < k_2 * \text{global variance}$, where $k_1 < k_2$.
Multiply identified pixels by constant $E > 1$. Leave other pixels alone.

Example: Local Statistics (3)

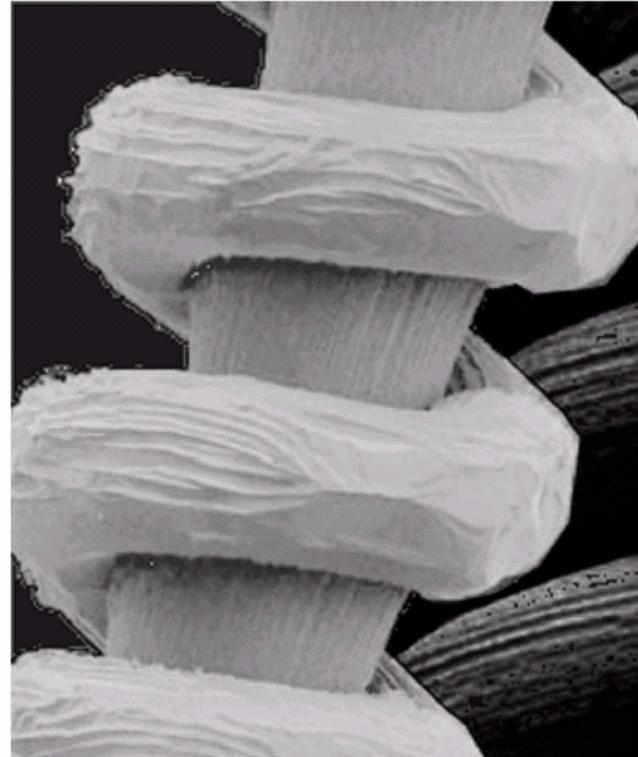


FIGURE 3.26
Enhanced SEM
image. Compare
with Fig. 3.24. Note
in particular the
enhanced area on
the right side of
the image.

Results for $E=4$, $k_0=0.4$, $k_1=0.02$, $k_2=0.4$. 3x3 neighborhoods used.