

---

# Digital Halftoning

Prof. George Wolberg  
Dept. of Computer Science  
City College of New York

# Objectives

---

- In this lecture we review digital halftoning techniques to convert grayscale images to bitmaps:
  - Unordered (random) dithering
  - Ordered dithering
  - Patterning
  - Error diffusion

# Background

---

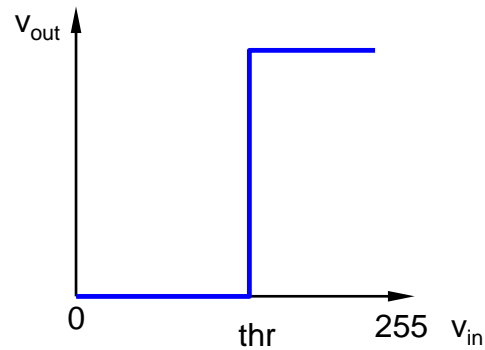
- An 8-bit grayscale image allows 256 distinct gray levels.
- Such images can be displayed on a computer monitor if the hardware supports the required number of intensity levels.
- However, some output devices print or display images with much fewer gray levels.
- In these cases, the grayscale images must be converted to binary images, where pixels are only black (0) or white (255).
- Thresholding is a poor choice due to objectionable artifacts.
- Strategy: sprinkle black-and-white dots to simulate gray.
- Exploit spatial integration (averaging) performed by eye.

# Thresholding

- The simplest way to convert from grayscale to binary.



8 bpp (256 levels)



1 bpp (two-level)

Loss of information is unacceptable.


# Unordered Dither (1)

---

- Reduce quantization error by adding uniformly distributed white noise (dither signal) to the input image prior to quantization.
- Dither hides objectional artifacts.
- To each pixel of the image, add a random number in the range  $[-m, m]$ , where  $m$  is  $\text{MXGRAY}/\text{quantization-levels}$ .




8 bpp (256 levels)



3 bpp (8 levels)

# Unordered Dither (2)


---



# Ordered Dithering

---

- Objective: expand the range of available intensities.
- Simulates  $n$  bpp images with  $m$  bpp, where  $n > m$  (usually  $m = 1$ ).
- Exploit eye's spatial integration.
  - Gray is due to average of black/white dot patterns.
  - Each dot is a circle of black ink whose area is proportional to  $(1 - \text{intensity})$ .
  - Graphics output devices approximate the variable circles of halftone reproductions.



- $2 \times 2$  pixel area of a bilevel display produces 5 intensity levels.
- $n \times n$  group of bilevel pixels produces  $n^2 + 1$  intensity levels.
- Tradeoff: spatial vs. intensity resolution.

# Dither Matrix (1)

---

- Consider the following 2x2 and 3x3 dither matrices:

$$D^{(2)} = \begin{bmatrix} 0 & 2 \\ 3 & 1 \end{bmatrix} \quad D^{(3)} = \begin{bmatrix} 6 & 8 & 4 \\ 1 & 0 & 3 \\ 5 & 2 & 7 \end{bmatrix}$$

- To display a pixel of intensity  $I$ , we turn on all pixels whose associated dither matrix values are less than  $I$ .
- The recurrence relation given below generates larger dither matrices of dimension  $n \times n$ , where  $n$  is a power of 2.


$$D^{(n)} = \begin{bmatrix} 4D^{(n/2)} + D_{00}^{(2)}U^{(n/2)} & 4D^{(n/2)} + D_{01}^{(2)}U^{(n/2)} \\ 4D^{(n/2)} + D_{10}^{(2)}U^{(n/2)} & 4D^{(n/2)} + D_{11}^{(2)}U^{(n/2)} \end{bmatrix}$$

where  $U^{(n)}$  is an  $n \times n$  matrix of 1's.




# Dither Matrix (2)

- Example: a 4x4 dither matrix can be derived from the 2x2 matrix.



# Patterning

- Let the output image be larger than the input image.
- Quantize the input image to  $[0 \dots n^2]$  gray levels.
- Threshold each pixel against all entries in the dither matrix.
  - Each pixel forms a  $4 \times 4$  block of black-and-white dots for a  $D^{(4)}$  matrix.
  - An  $n \times n$  input image becomes a  $4n \times 4n$  output image.
- Multiple display pixels per input pixel.
- The dither matrix  $D_{ij}^{(n)}$  is used as a spatially-varying threshold.
- Large input areas of constant value are displayed exactly as before.



# Implementation

---

- Let the input and output images share the same size.
- First quantize the input image to  $[0 \dots n^2]$  gray levels.
- Compare the dither matrix with the input image.

```
for(y=0; y<h; y++)          // visit all input rows
    for(x=0; x<w; x++){      // visit all input cols
        i = x % n;           // dither matrix index
        j = y % n;           // dither matrix index

        // threshold pixel using dither value  $D_{ij}^{(n)}$ 
        out[y*w+x] = (in[y*w+x] >  $D_{ij}^{(n)}$ ) ? 255 : 0;
    }
```

# Examples

---



**8 bpp (256 levels)**



**1 bpp ( $D^3$ )**



**1 bpp ( $D^4$ )**



**1 bpp ( $D^8$ )**

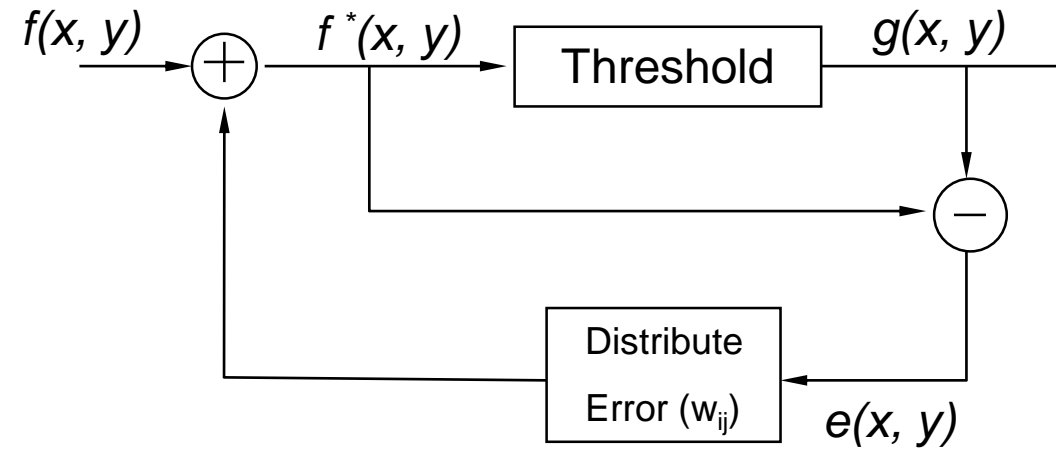
# Error Diffusion

---

- An error is made every time a grayvalue is assigned to be black or white at the output.
- Spread that error to its neighbors to compensate for over/undershoots in the output assignments
  - If input pixel 130 is mapped to white (255) then its excessive brightness ( $255 - 130$ ) must be subtracted from neighbors to enforce a bias towards darker values to compensate for the excessive brightness.
- Like ordered dithering, error diffusion permits the output image to share the same dimension as the input image.

# Floyd-Steinberg Algorithm

---



$$f^*(x, y) = f(x, y) + \sum_i \sum_j w_{ij} e(x-i, y-j) = \text{"corrected intensity value"}$$

$$g(x, y) = \begin{cases} 255 & \text{if } f^*(x, y) > MXGRAY / 2 \\ 0 & \text{otherwise} \end{cases}$$


$$e(x, y) = f^*(x, y) - g(x, y)$$

$$\sum_i \sum_j w_{ij} = 1$$


# Error Diffusion Weights

---


- Note that visual improvements are possible if left-to-right scanning among rows is replaced by serpentine scanning (zig-zag). That is, scan odd rows from left-to right, and scan even rows from right-to-left.
- Further improvements can be made by using larger neighborhoods.
- The sum of the weights should equal 1 to avoid emphasizing or suppressing the spread of errors.



**Floyd-Steinberg**



**Jarvis-Judice-Ninke**



**Stucki**



# Examples (1)

---



**Floyd-Steinberg**



**Jarvis-Judice-Ninke**






## Examples (2)


---



**Floyd-Steinberg**



**Jarvis-Judice-Ninke**






# Examples (3)


---



**Floyd-Steinberg**



**Jarvis-Judice-Ninke**



# Implementation

---

```
thr = MXGRAY /2;                // init threshold value
for(y=0; y<h; y++){              // visit all input rows
    for(x=0; x<w; x++) {          // visit all input cols
        *out = (*in < thr)?        // threshold
            BLACK : WHITE;         // note: use LUT!

        e = *in - *out;            // eval error
        in[ 1 ] +=(e*7/16.);        // add error to E  nbr
        in[w-1] +=(e*3/16.);        // add error to SW nbr
        in[ w ] +=(e*5/16.);        // add error to S  nbr
        in[w+1] +=(e*1/16.);        // add error to SE nbr


        in++;                      // advance input  ptr
        out++;                     // advance output ptr
    }
}
```

# Comments

- Two potential problems complicate implementation:

- errors can be deposited beyond image border
- errors may force pixel grayvalues outside the  $[0,255]$  range


← True for all neighborhood ops



# Solutions to Border Problem (1)

---

- Perform **if** statement prior to every error deposit
  - Drawback: inefficient / slow
- Limit excursions of sliding weights to lie no closer than 1 pixel from image boundary (2 pixels for J-J-N weights).
  - Drawback: output will be smaller than input
- Pad image with extra rows and columns so that limited excursions will yield smaller image that conforms with original input dimensions. Padding serves as placeholder.
  - Drawback: excessive memory needs for intermediate image




## Solutions to Border Problem (2)

---

- Use of padding is further undermined by fact that 16-bit precision (**short**) is needed to accommodate pixel values outside  $[0, 255]$  range.
- A better solution is suggested by fact that only two rows are active while processing a single scanline in the Floyd-Steinberg algorithm (3 for JJN).
- Therefore, use a 2-row (or 3-row) circular buffer to handle the two (or three) current rows.
- The circular buffer will have the necessary padding and 16-bit precision.
- This significantly reduces memory requirements.

# Circular Buffer

---



# New Implementation

---

```
thr = MXGRAY /2;           // init threshold value
copyRowToCircBuffer(0);     // copy row 0 to circular buffer
for(y=0; y<h; y++){        // visit all input rows
    copyRowToCircBuffer(y+1); // copy next row to circ buffer
    in1 = buf[ y %2] + 1;    // circ buffer ptr; skip over pad
    in2 = buf[(y+1)%2] + 1;  // circ buffer ptr; skip over pad
    for(x=0; x<w; x++) {    // visit all input cols
        *out = (*in1 < thr)? BLACK : WHITE; // threshold

        e = *in1 - *out;    // eval error
        in1[ 1] +=(e*7/16.); // add error to E  nbr
        in2[-1] +=(e*3/16.); // add error to SW nbr
        in2[ 0] +=(e*5/16.); // add error to S  nbr
        in2[ 1] +=(e*1/16.); // add error to SE nbr

        in1++; in2++;       // advance circ buffer ptrs
        out++;              // advance output ptr
    }
}
```